Apache Hadoop Framework to Process Big Data

Hadoop is a framework that allows us to store and process large data sets in parallel and distributed fashion

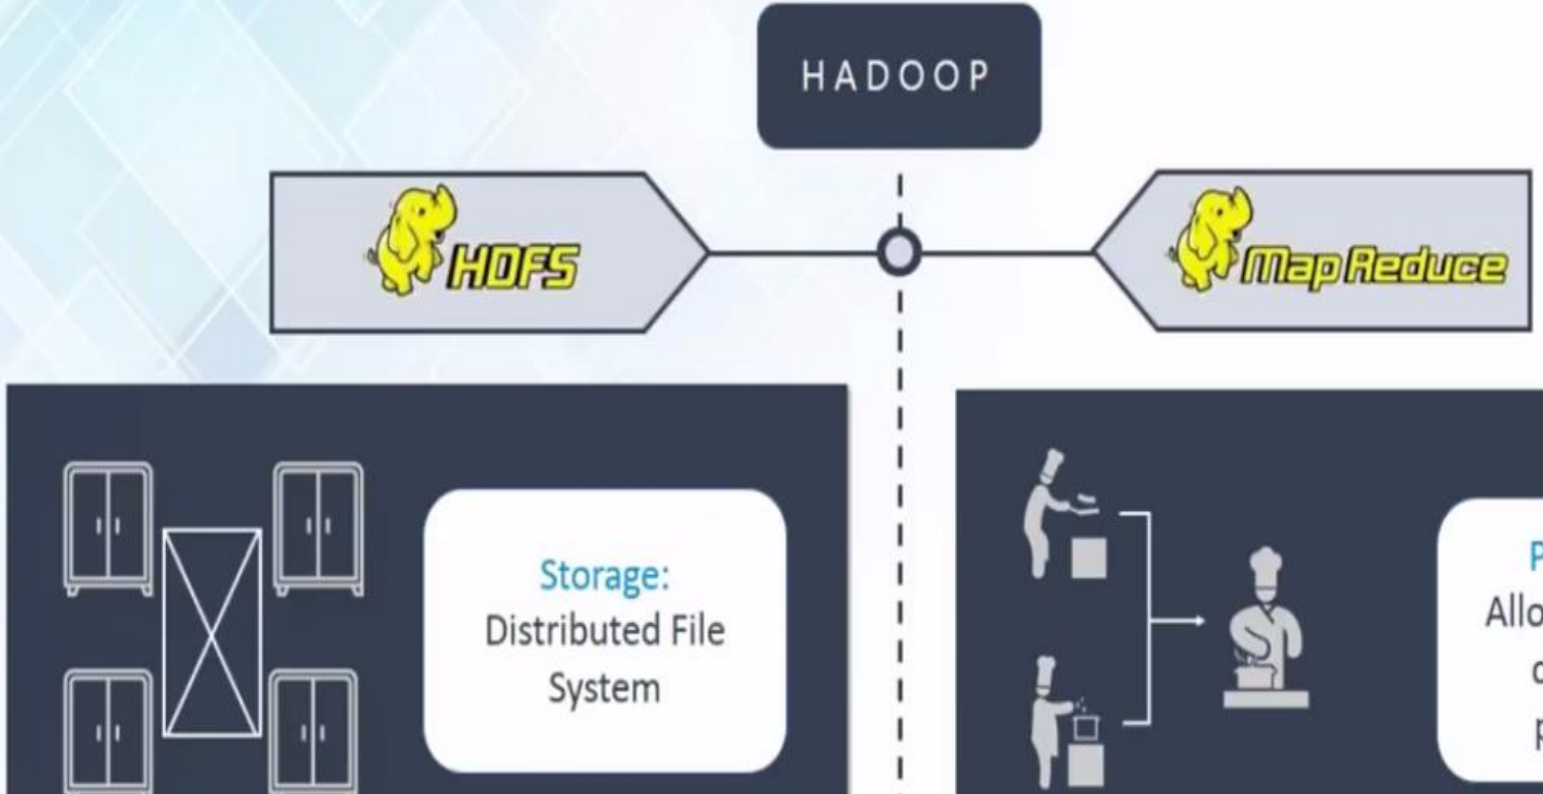

HADOOP CORE COMPONENTS

HDFS

MapReduce

Storage:
Distributed File System

Processing:
Allows parallel & distributed processing

HDFS Core Components

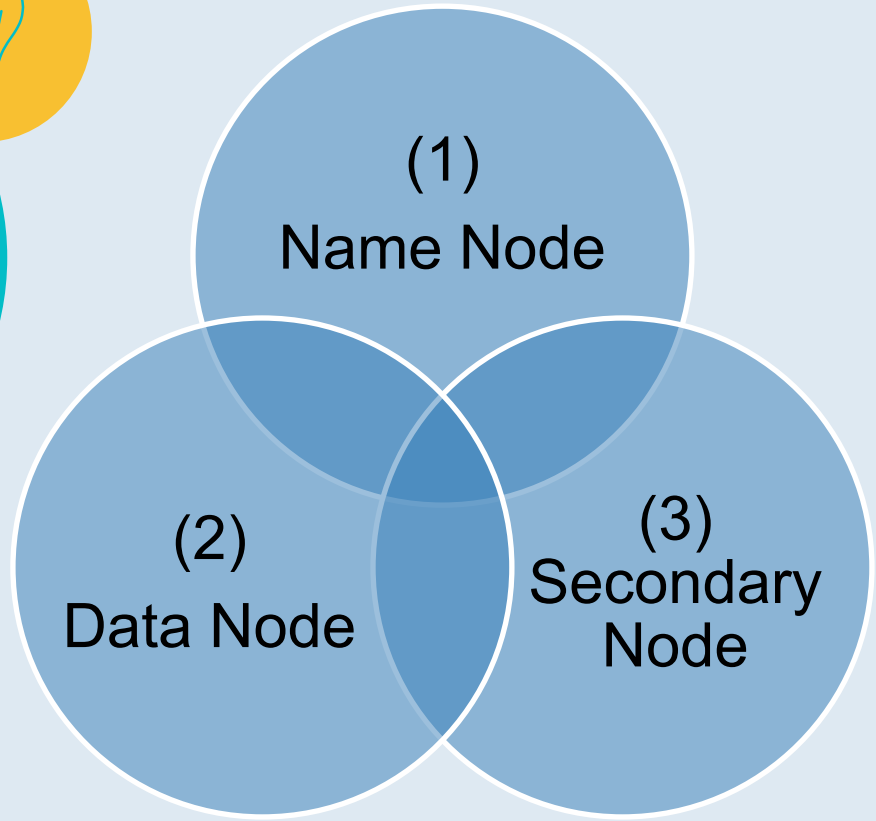

HDFS Core Components-Name Node & Data Node

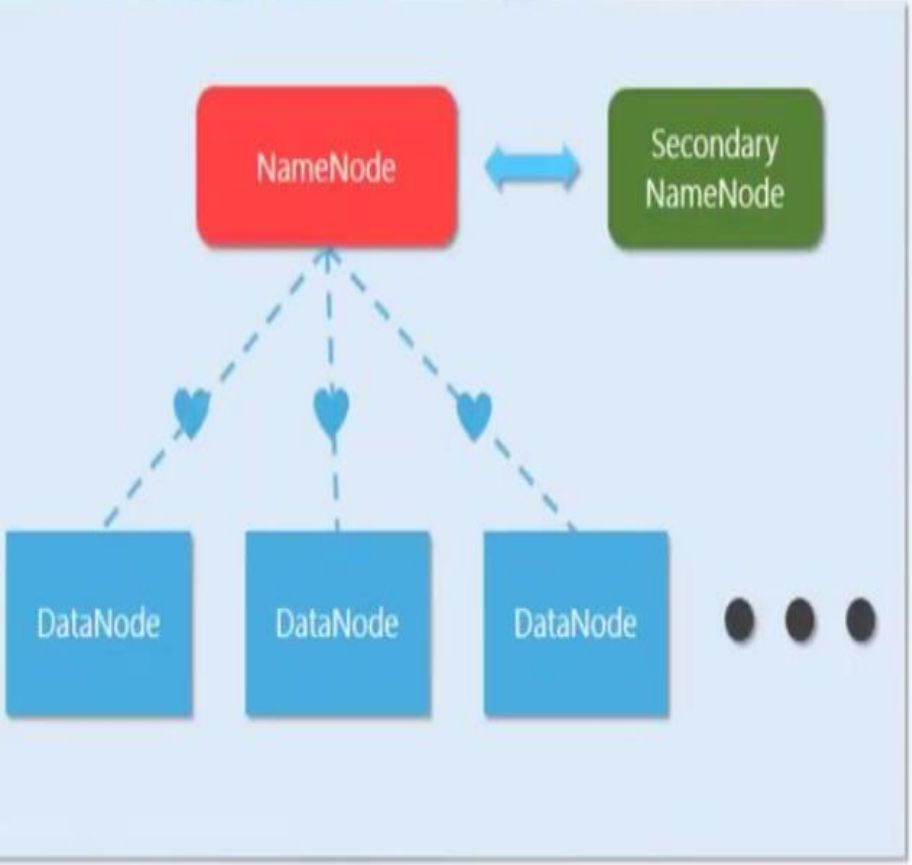

NameNode:

- Maintains and Manages DataNodes
- Records metadata i.e. information about data blocks e.g. location of blocks stored, the size of the files, permissions, hierarchy, etc.
- Receives heartbeat and block report from all the DataNodes

DataNode:

- Slave daemons
- Stores actual data
- Serves read and write requests from the clients

Secondary Name Node & Check pointing

- Checkpointing is a process of combining edit logs with FsImage
- Secondary NameNode takes over the responsibility of checkpointing, therefore, making NameNode more available
- Allows faster Failover as it prevents edit logs from getting too huge
- Checkpointing happens periodically (default: 1 hour)

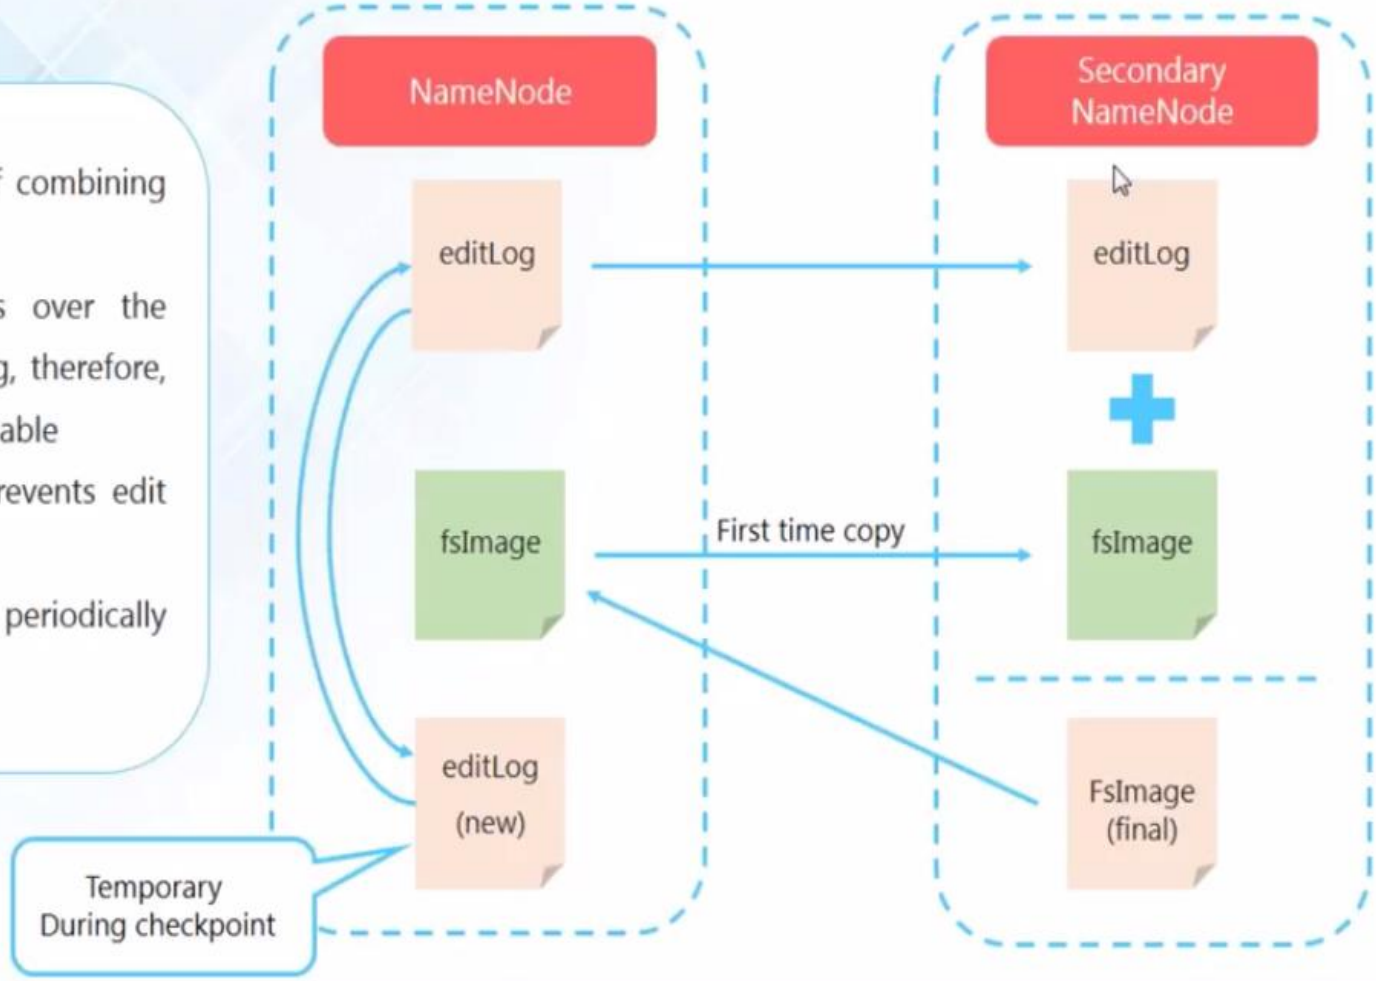

HDFS Data Blocks

- Each file is stored on HDFS as blocks
- The default size of each block is 128 MB in Apache Hadoop 2.x (64 MB in Apache Hadoop 1.x)

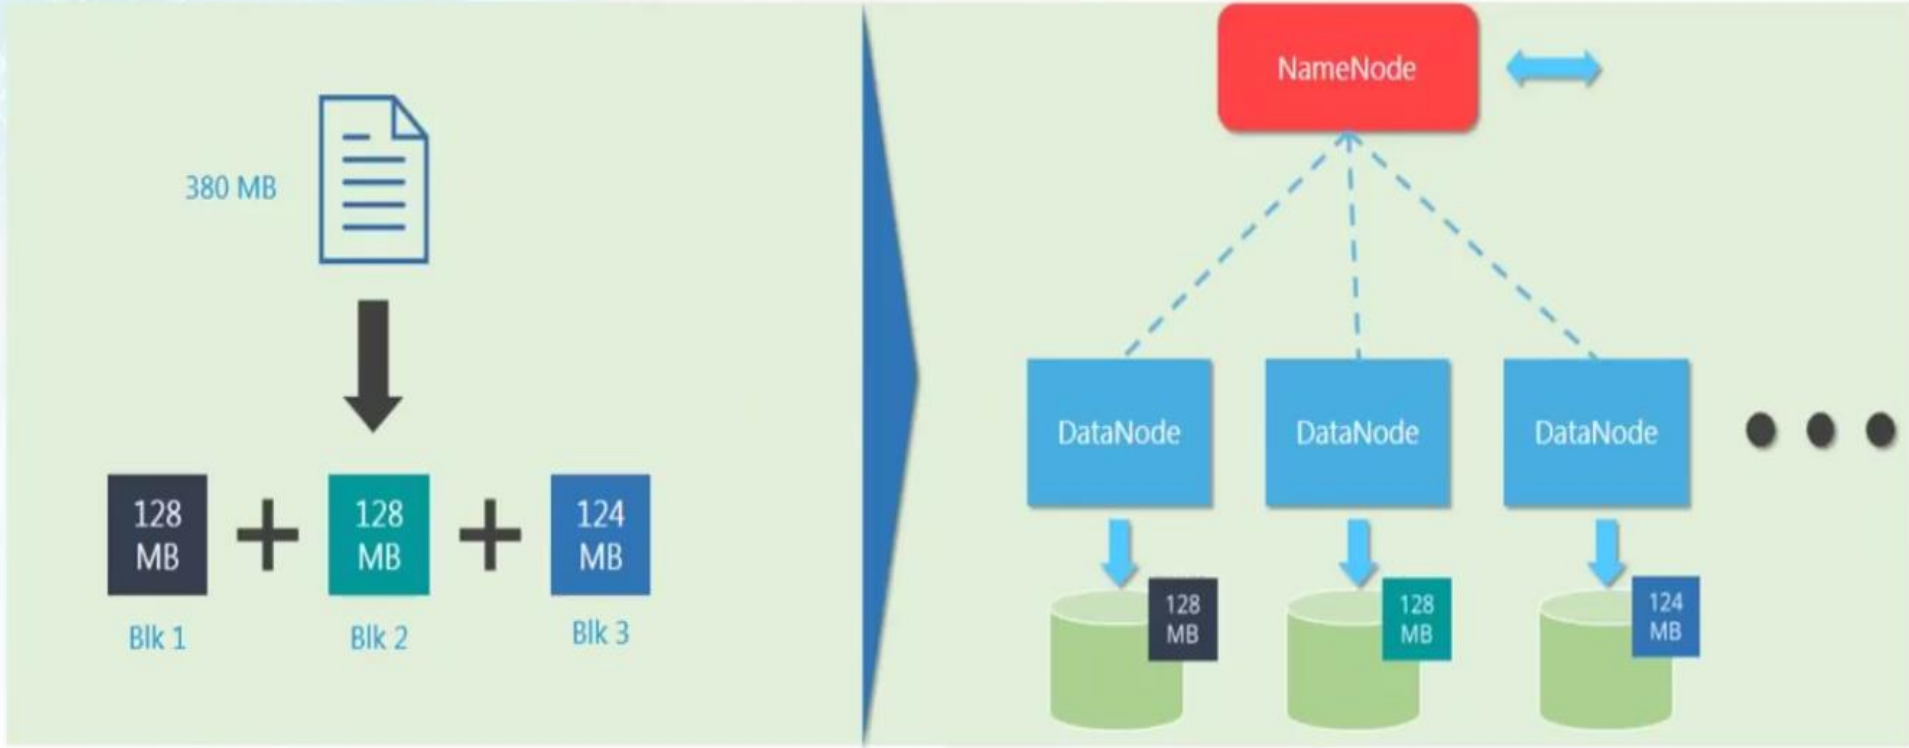

Fault Tolerance :How Hadoop cope up with DataNode Failure

Scenario:

One of the DataNodes crashed containing the data blocks

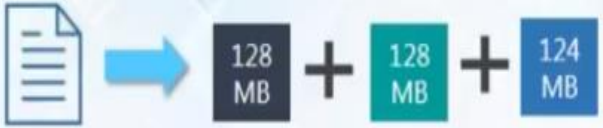

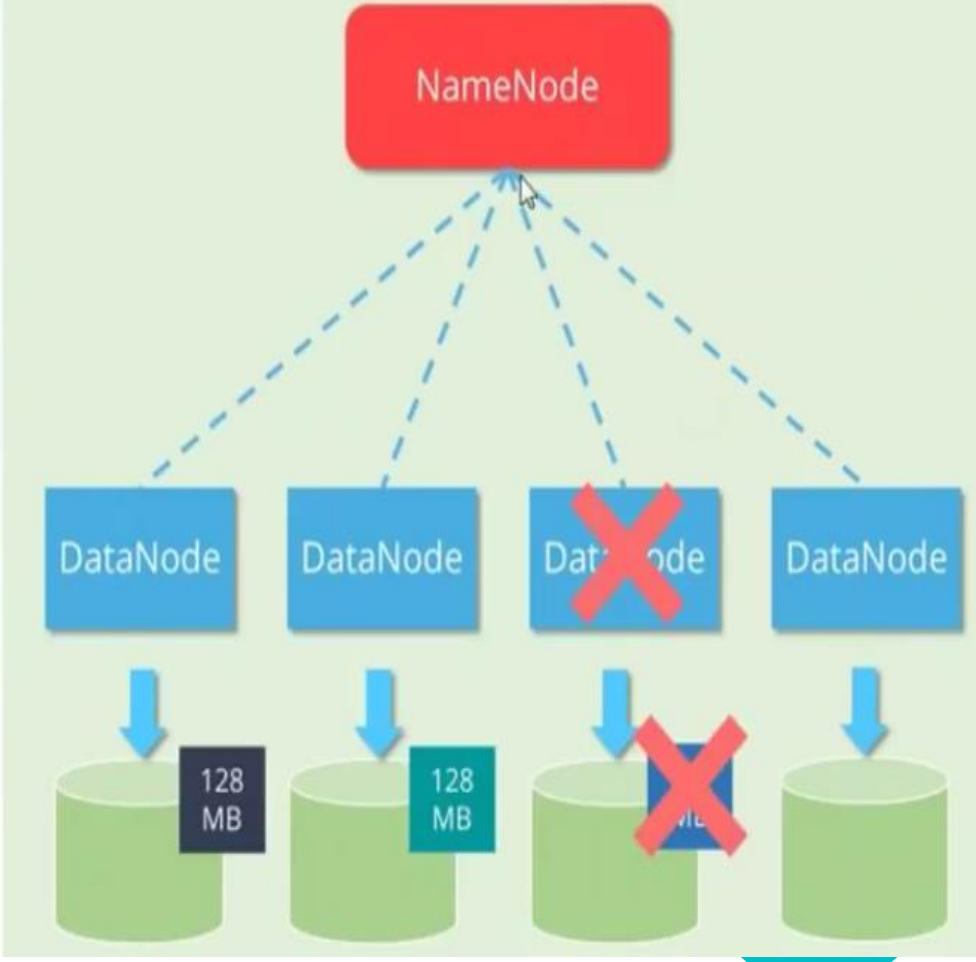

Fault Tolerance :Replication Factor

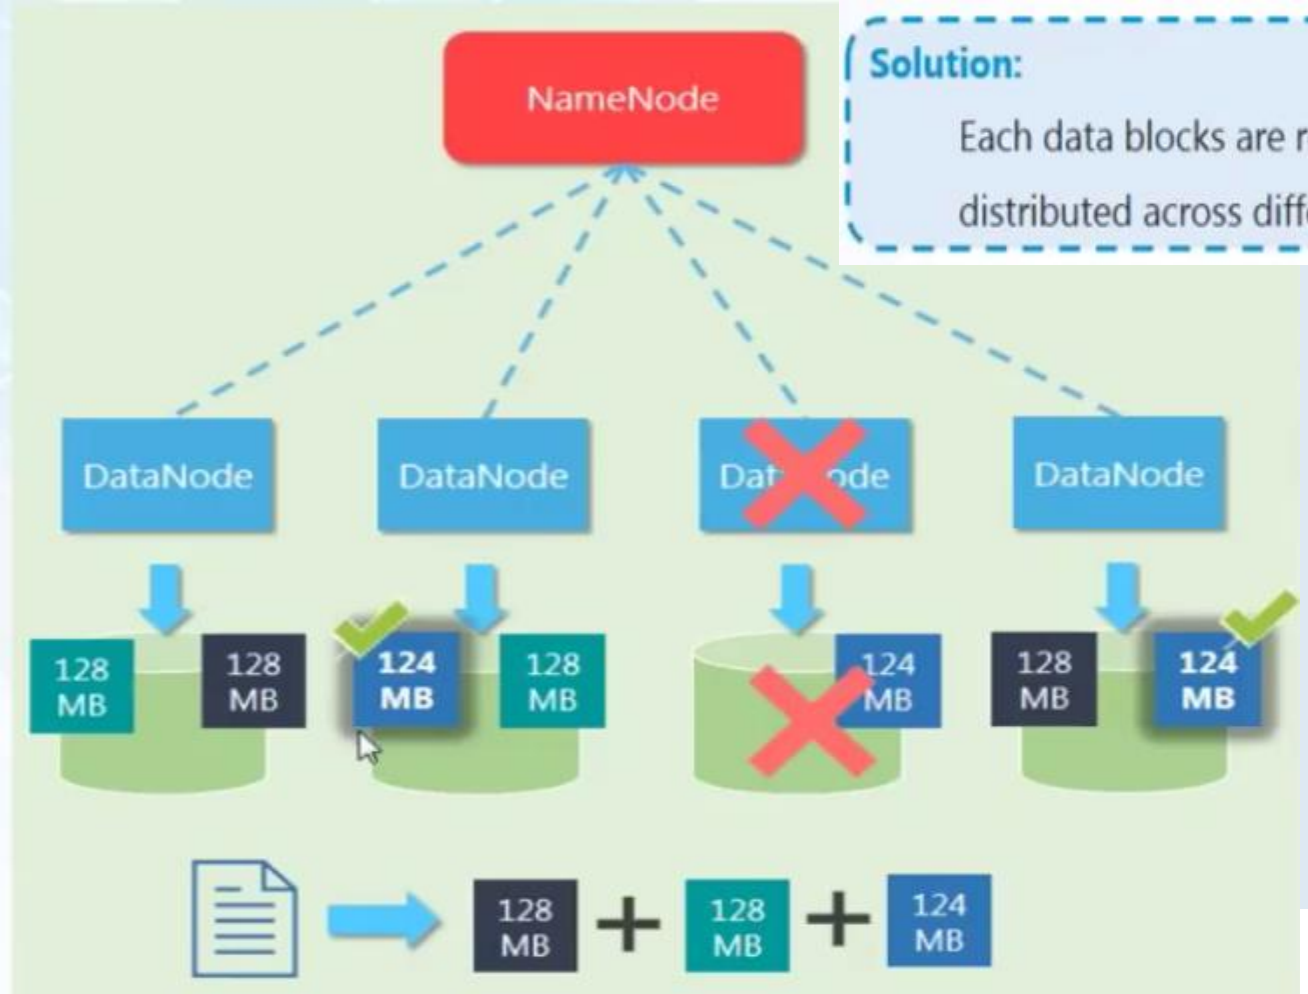

Solution:

Each data blocks are replicated (thrice by default) and are distributed across different DataNodes

As it is said Never Put All Your Eggs in the Same Basket

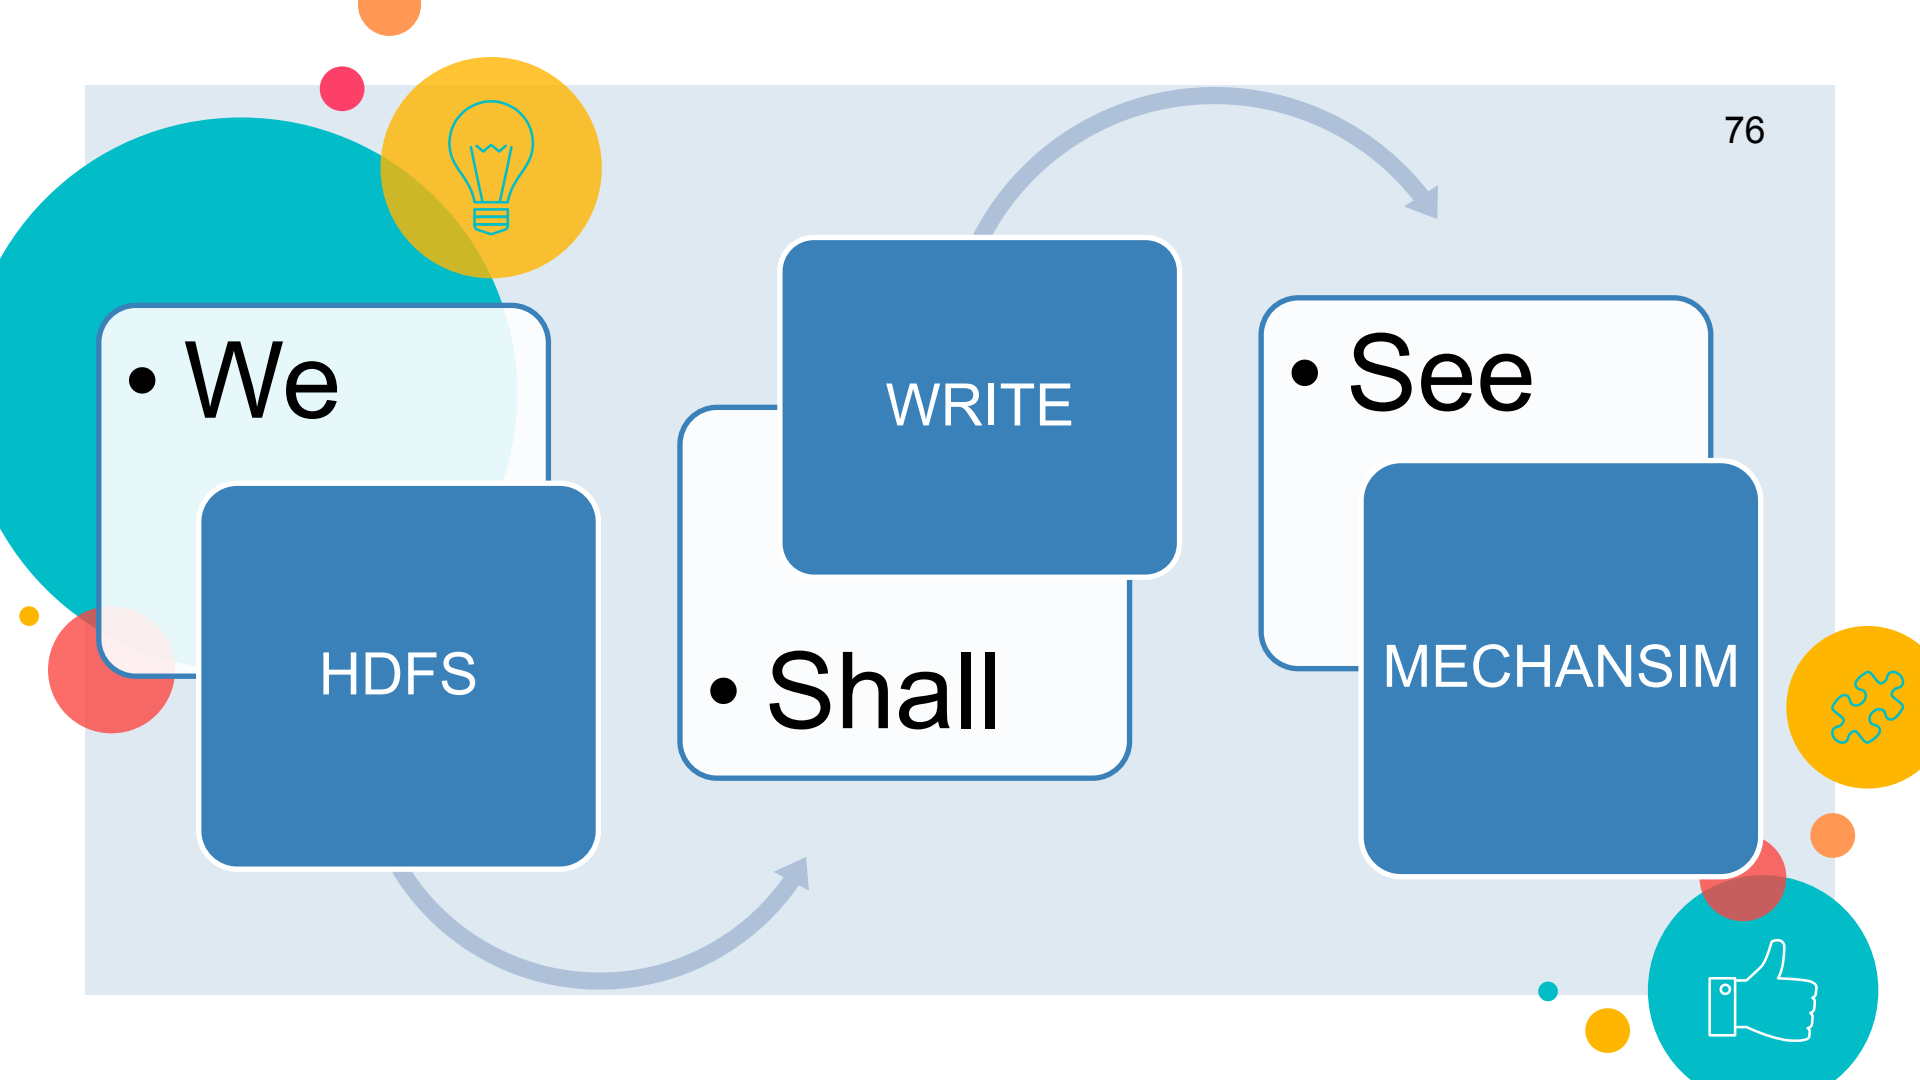
• We
HDFS
WRITE
• Shall
• See
MECHANSIM

HDFS Write Mechanism – Pipeline Setup

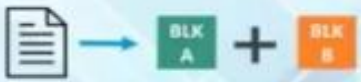

Setting up HDFS - Write Pipeline

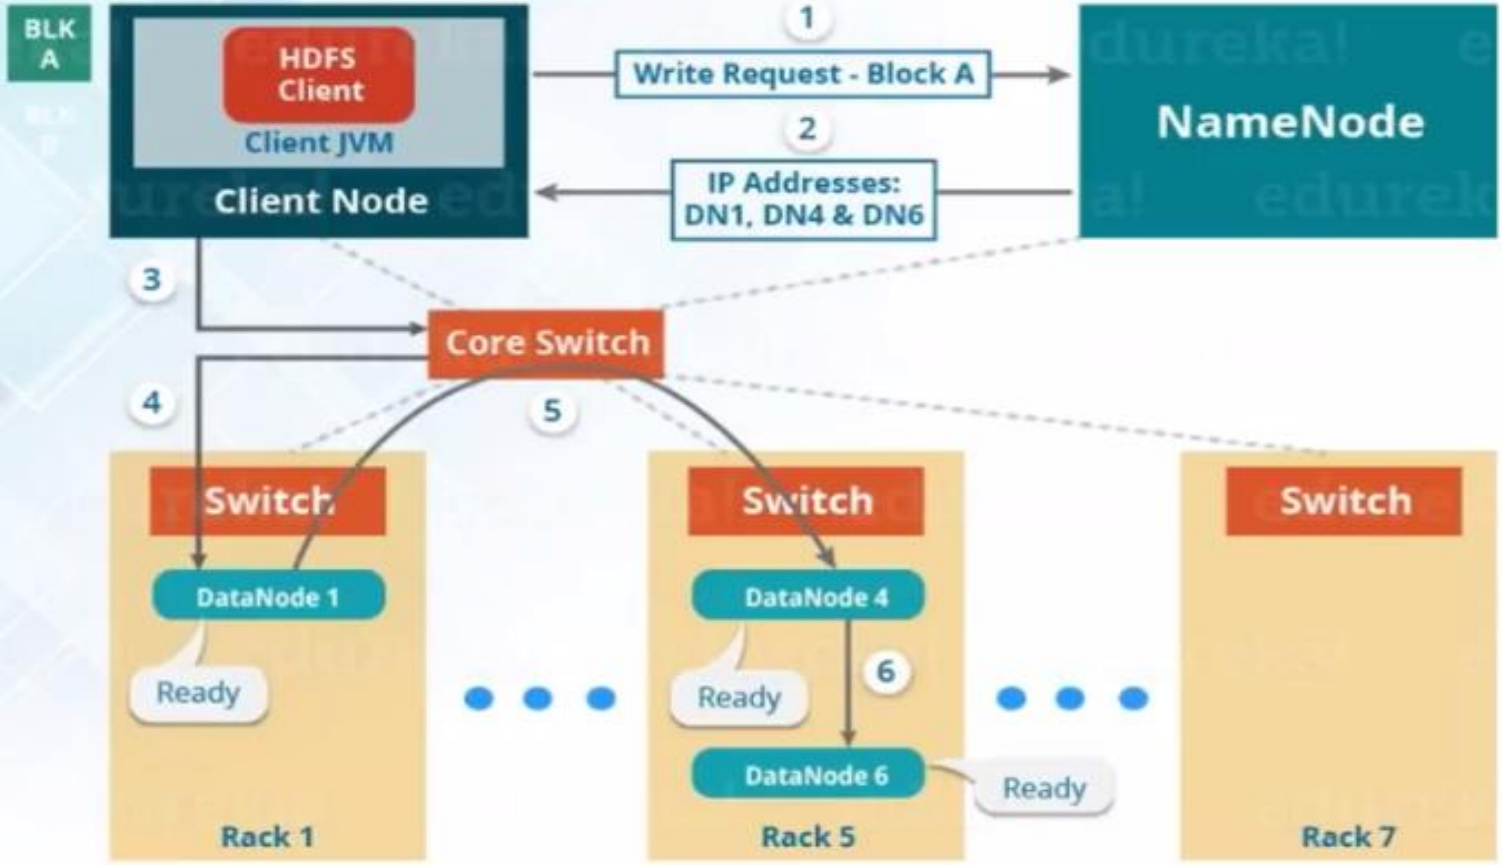

HDFS Write Mechanism – Writing a Block

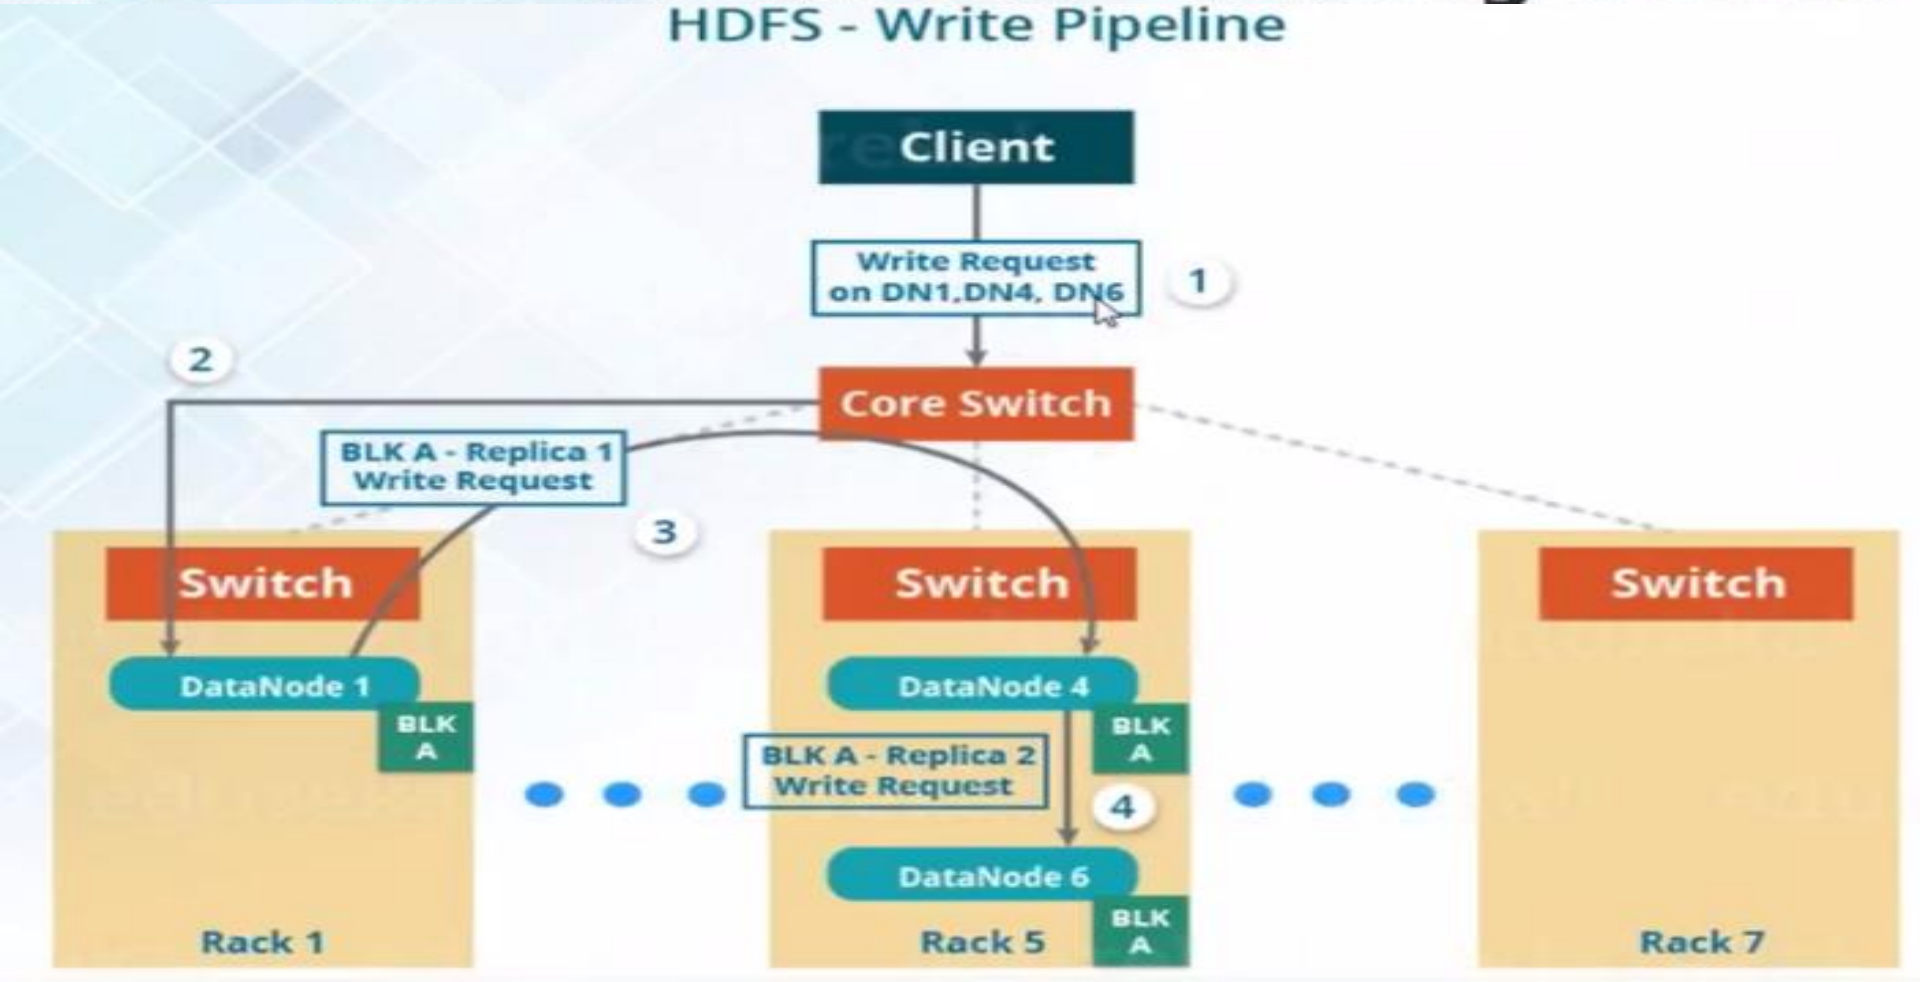

HDFS Write Mechanism - Acknowledgement

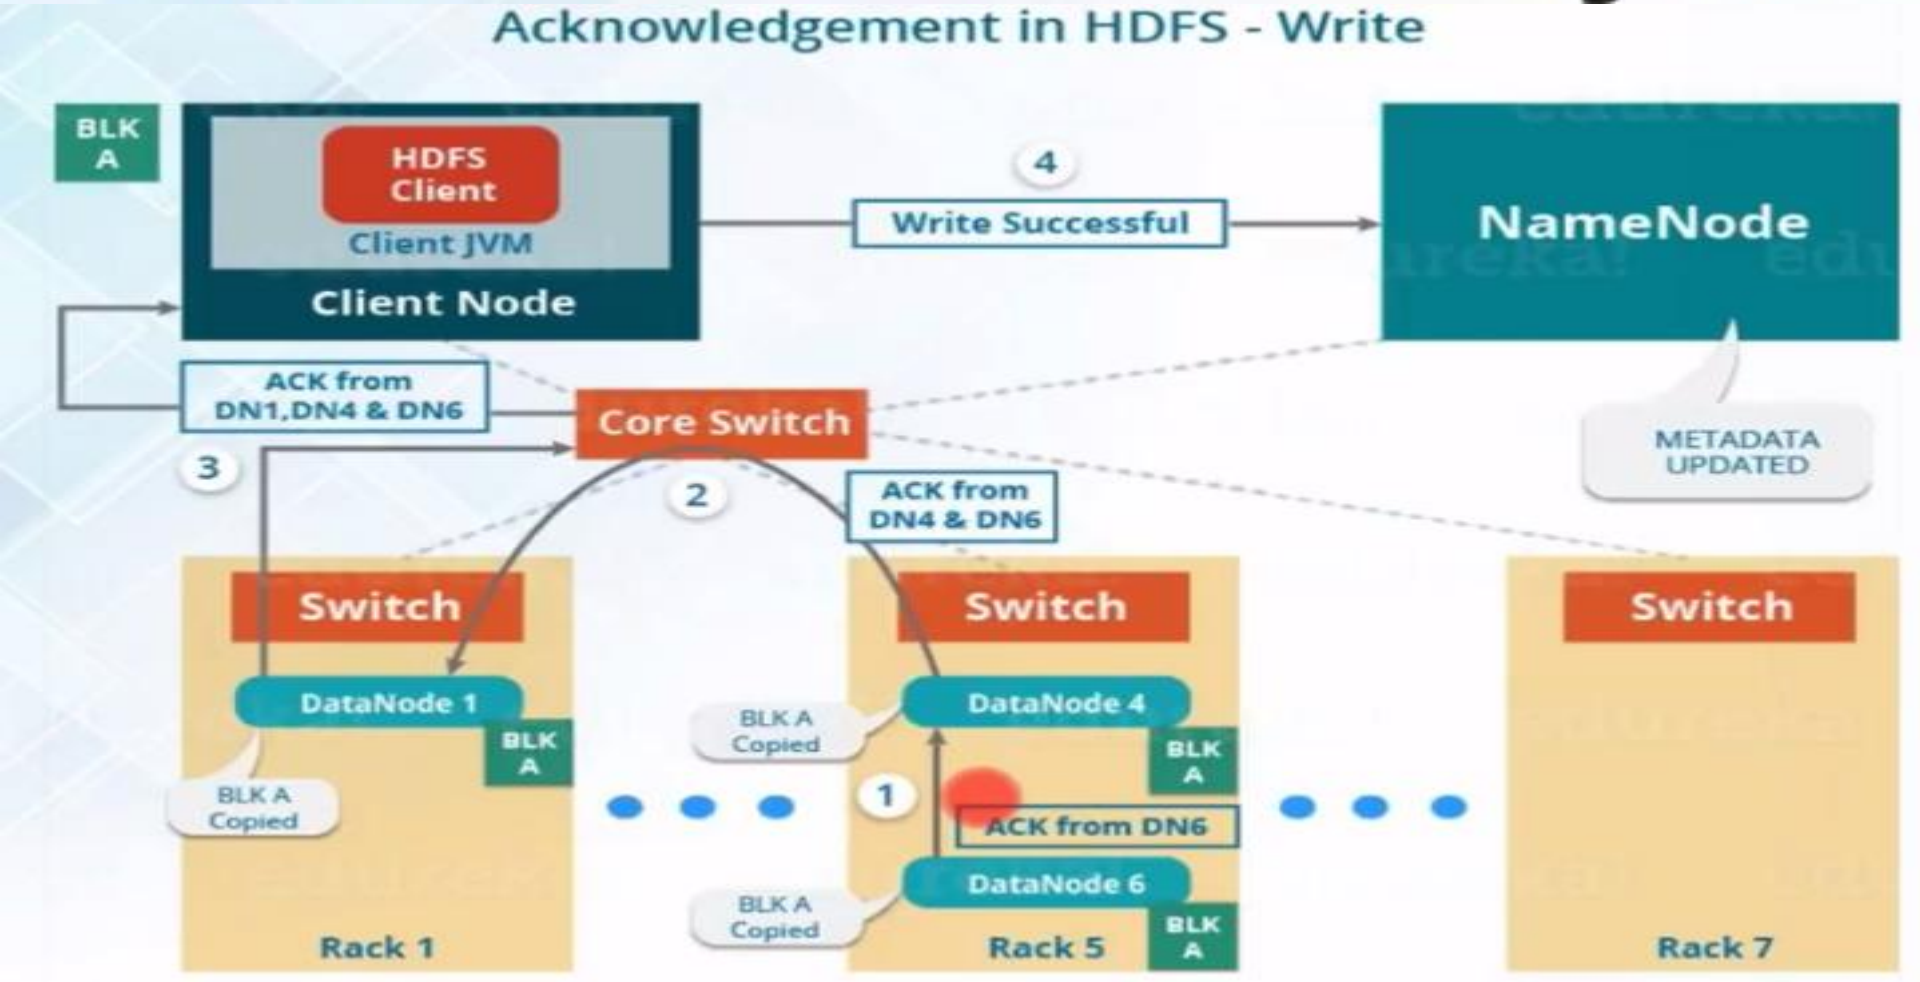

HDFS Multi-Block Write Mechanism

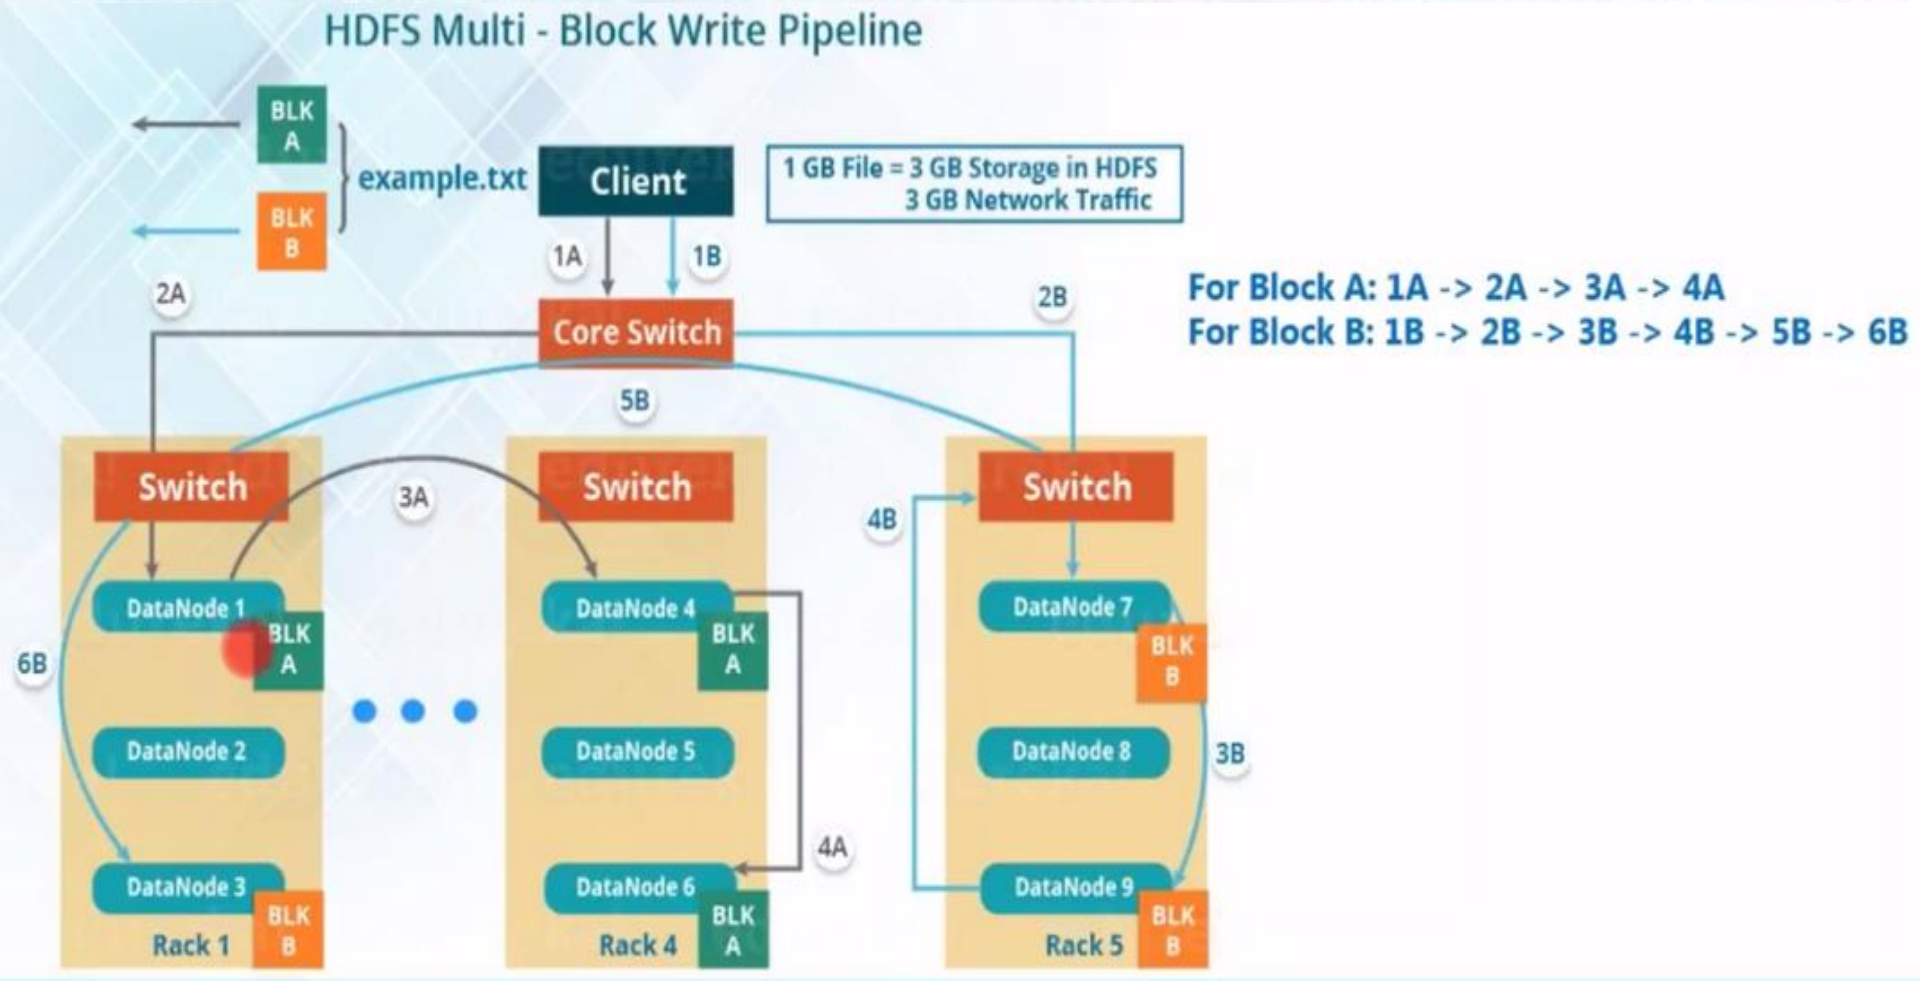

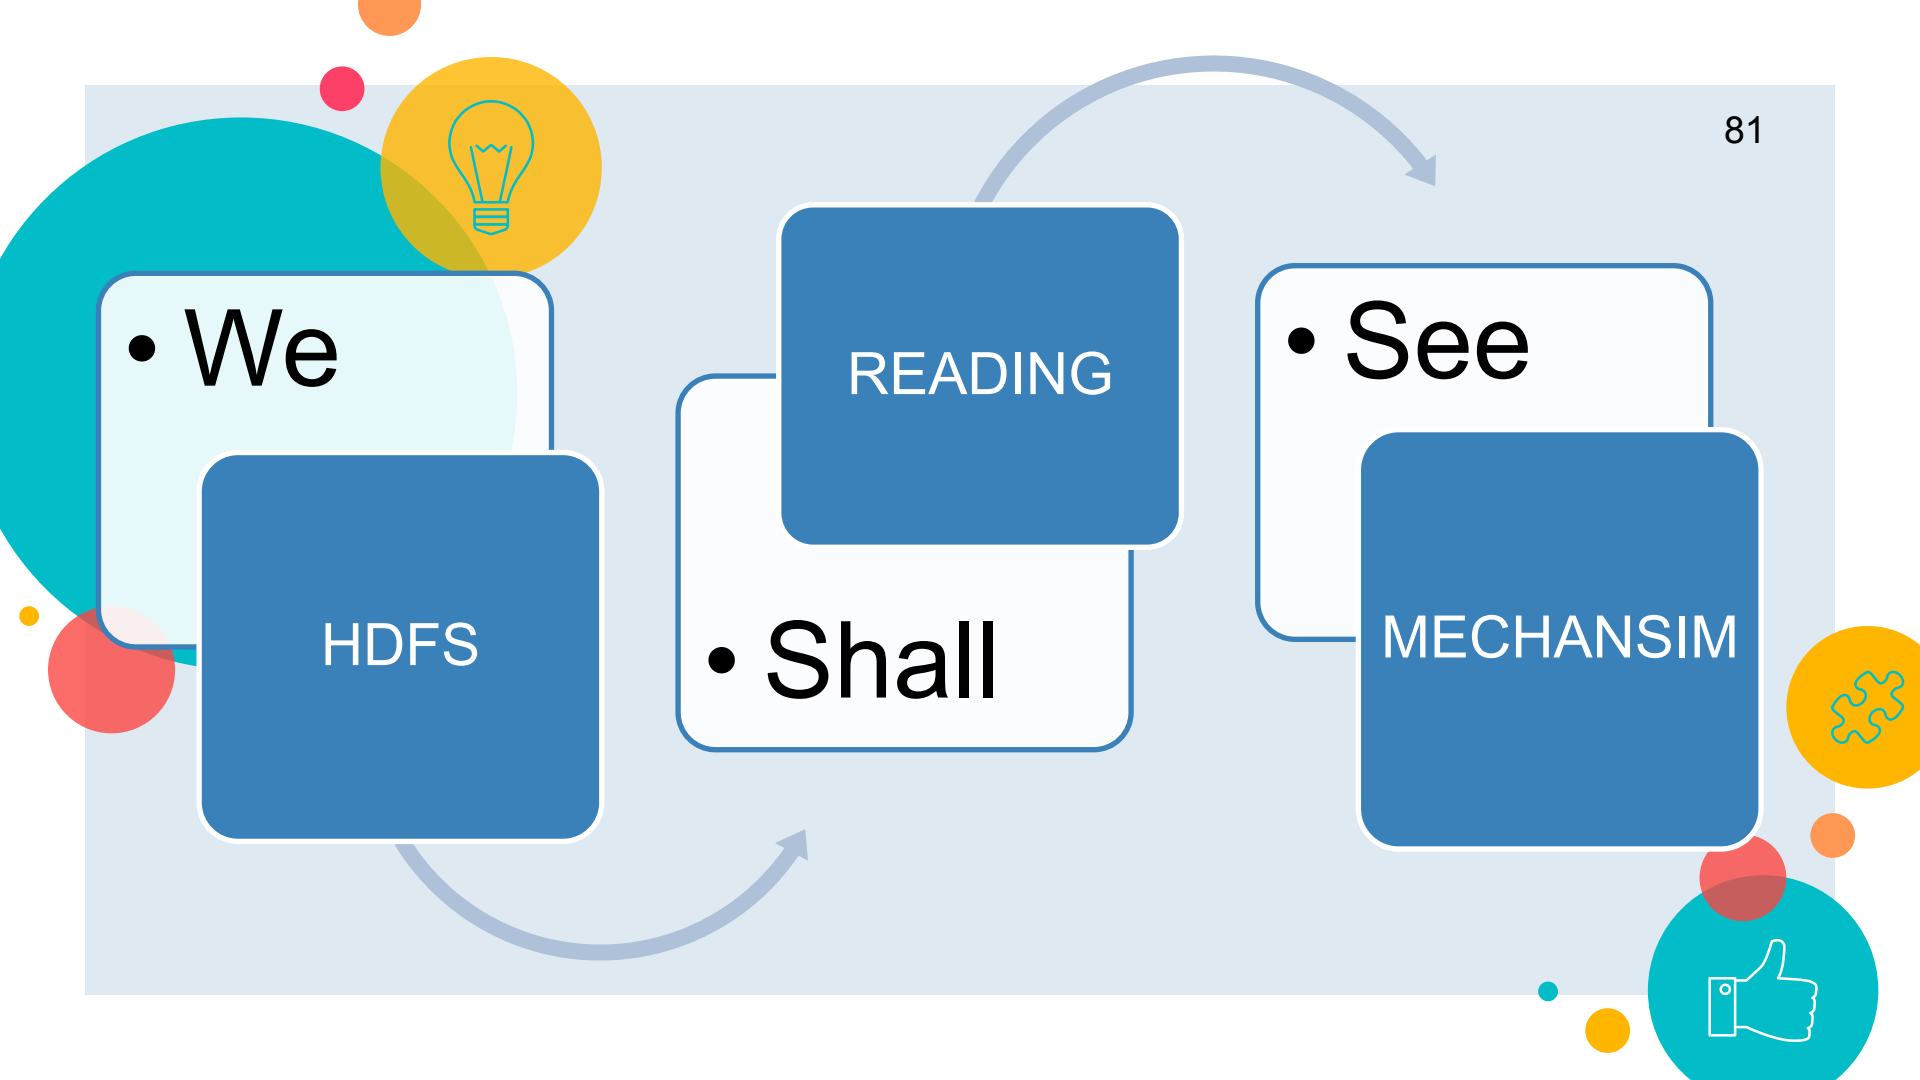
• We
HDFS
READING
• Shall
• See
MECHANSIM

HDFS Read Mechanism

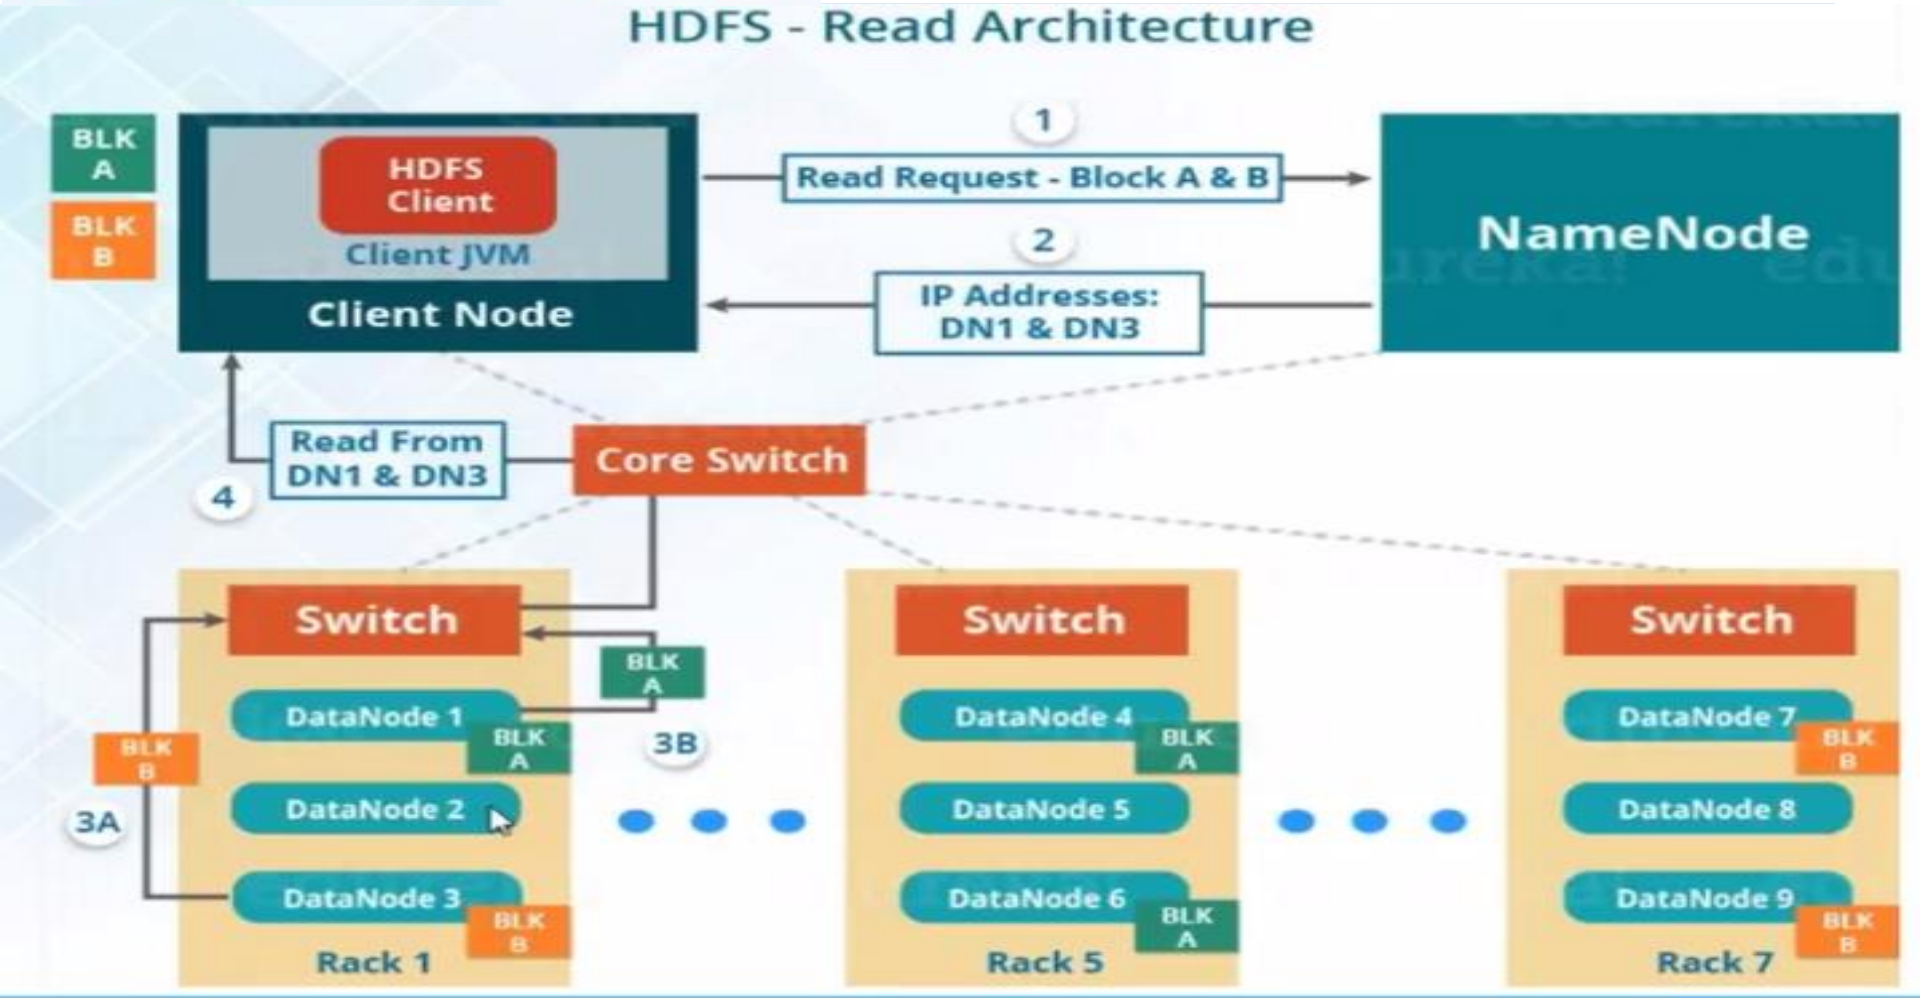

HADOOP CORE COMPONENTS

HDFS

MapReduce

Storage:
Distributed File System

Processing:
Allows parallel & distributed processing

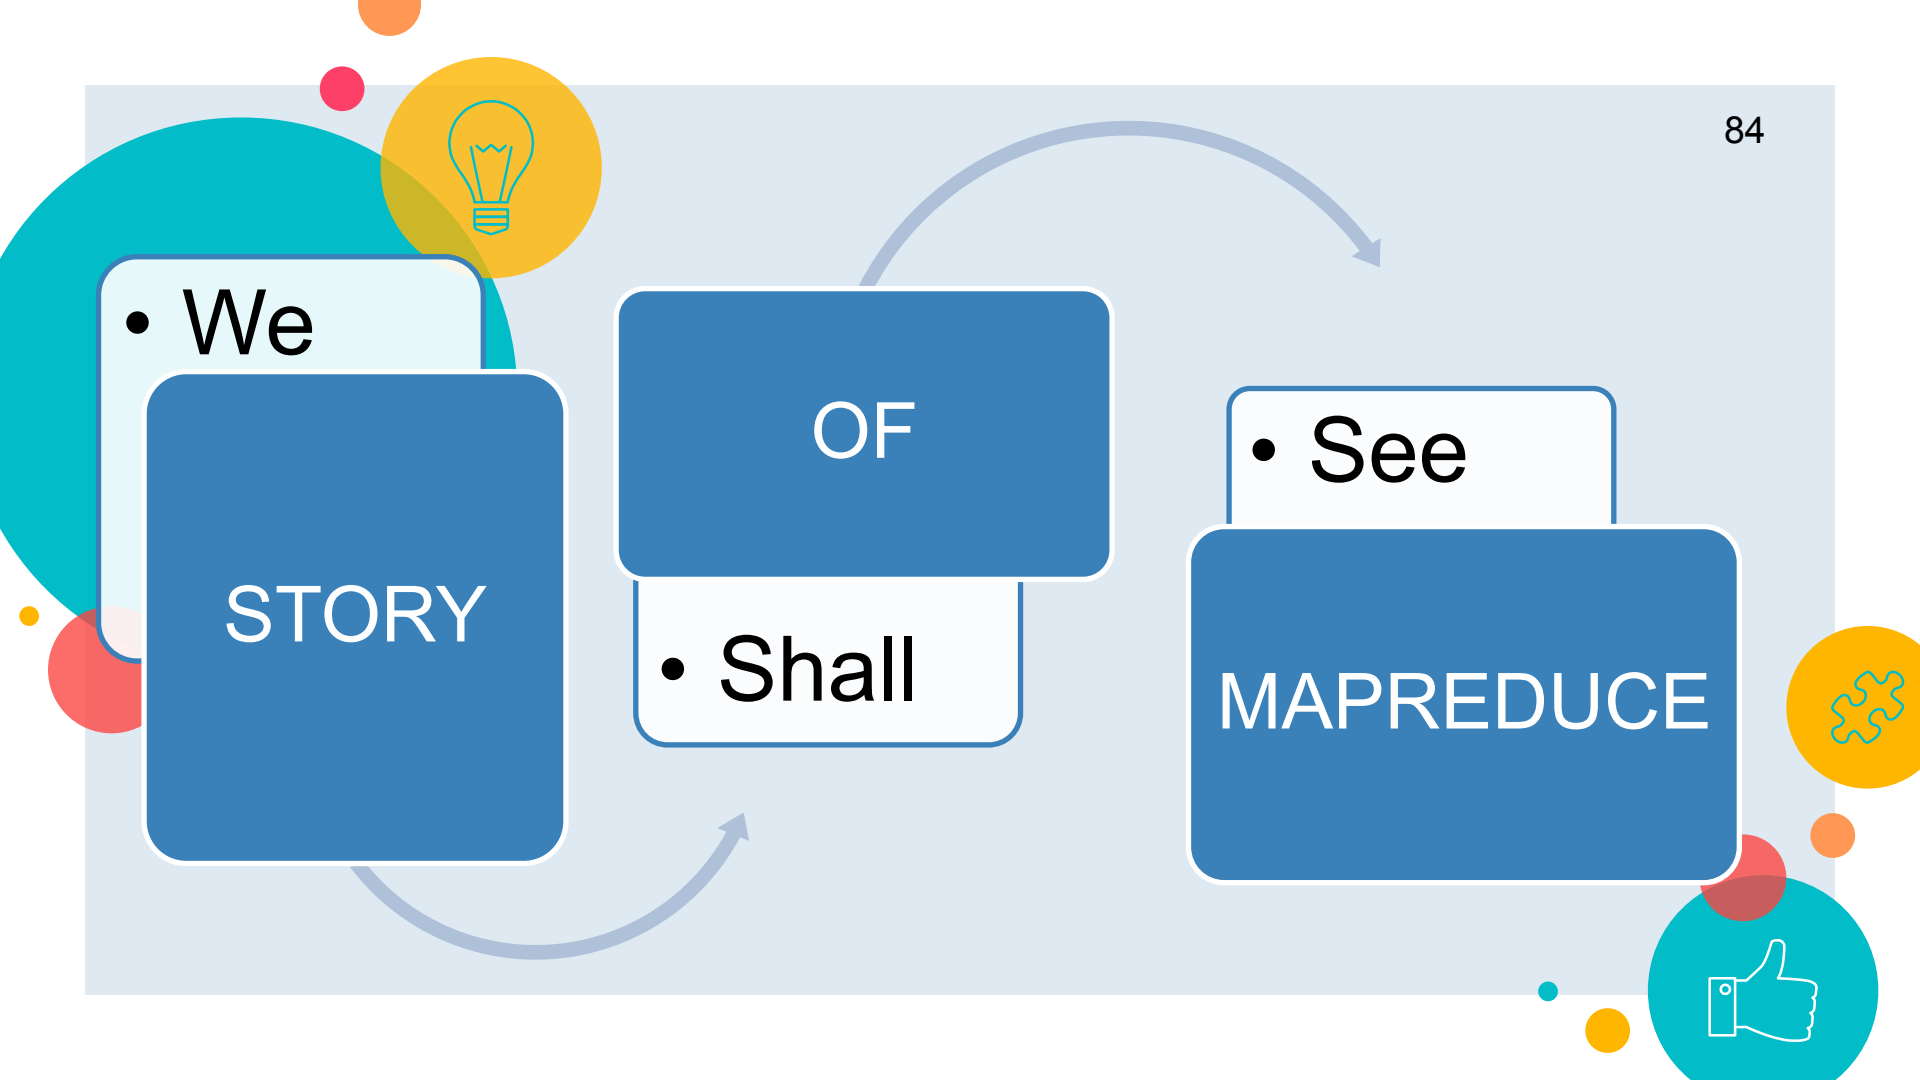
• We
STORY
OF
• Shall
• See
MAPREDUCE

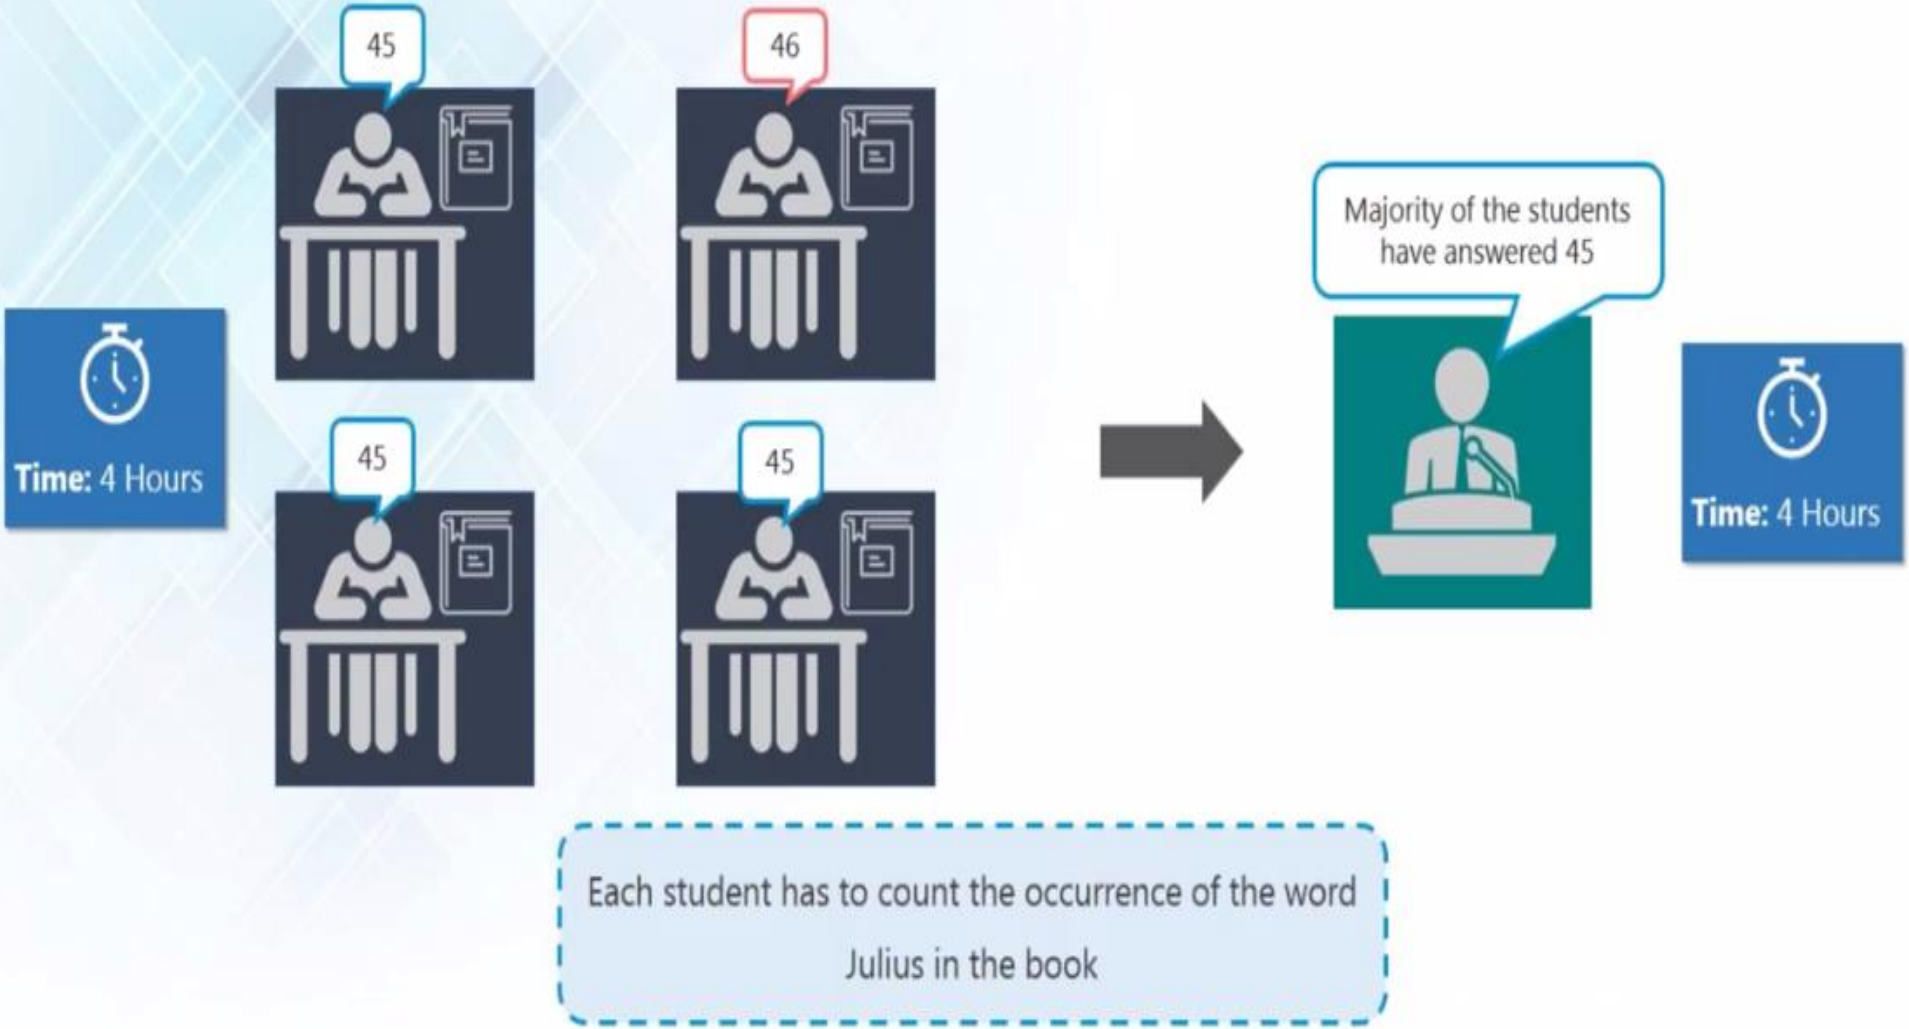
45
46
45
45
Time: 4 Hours
Majority of the students have answered 45
Time: 4 Hours
Each student has to count the occurrence of the word Julius in the book

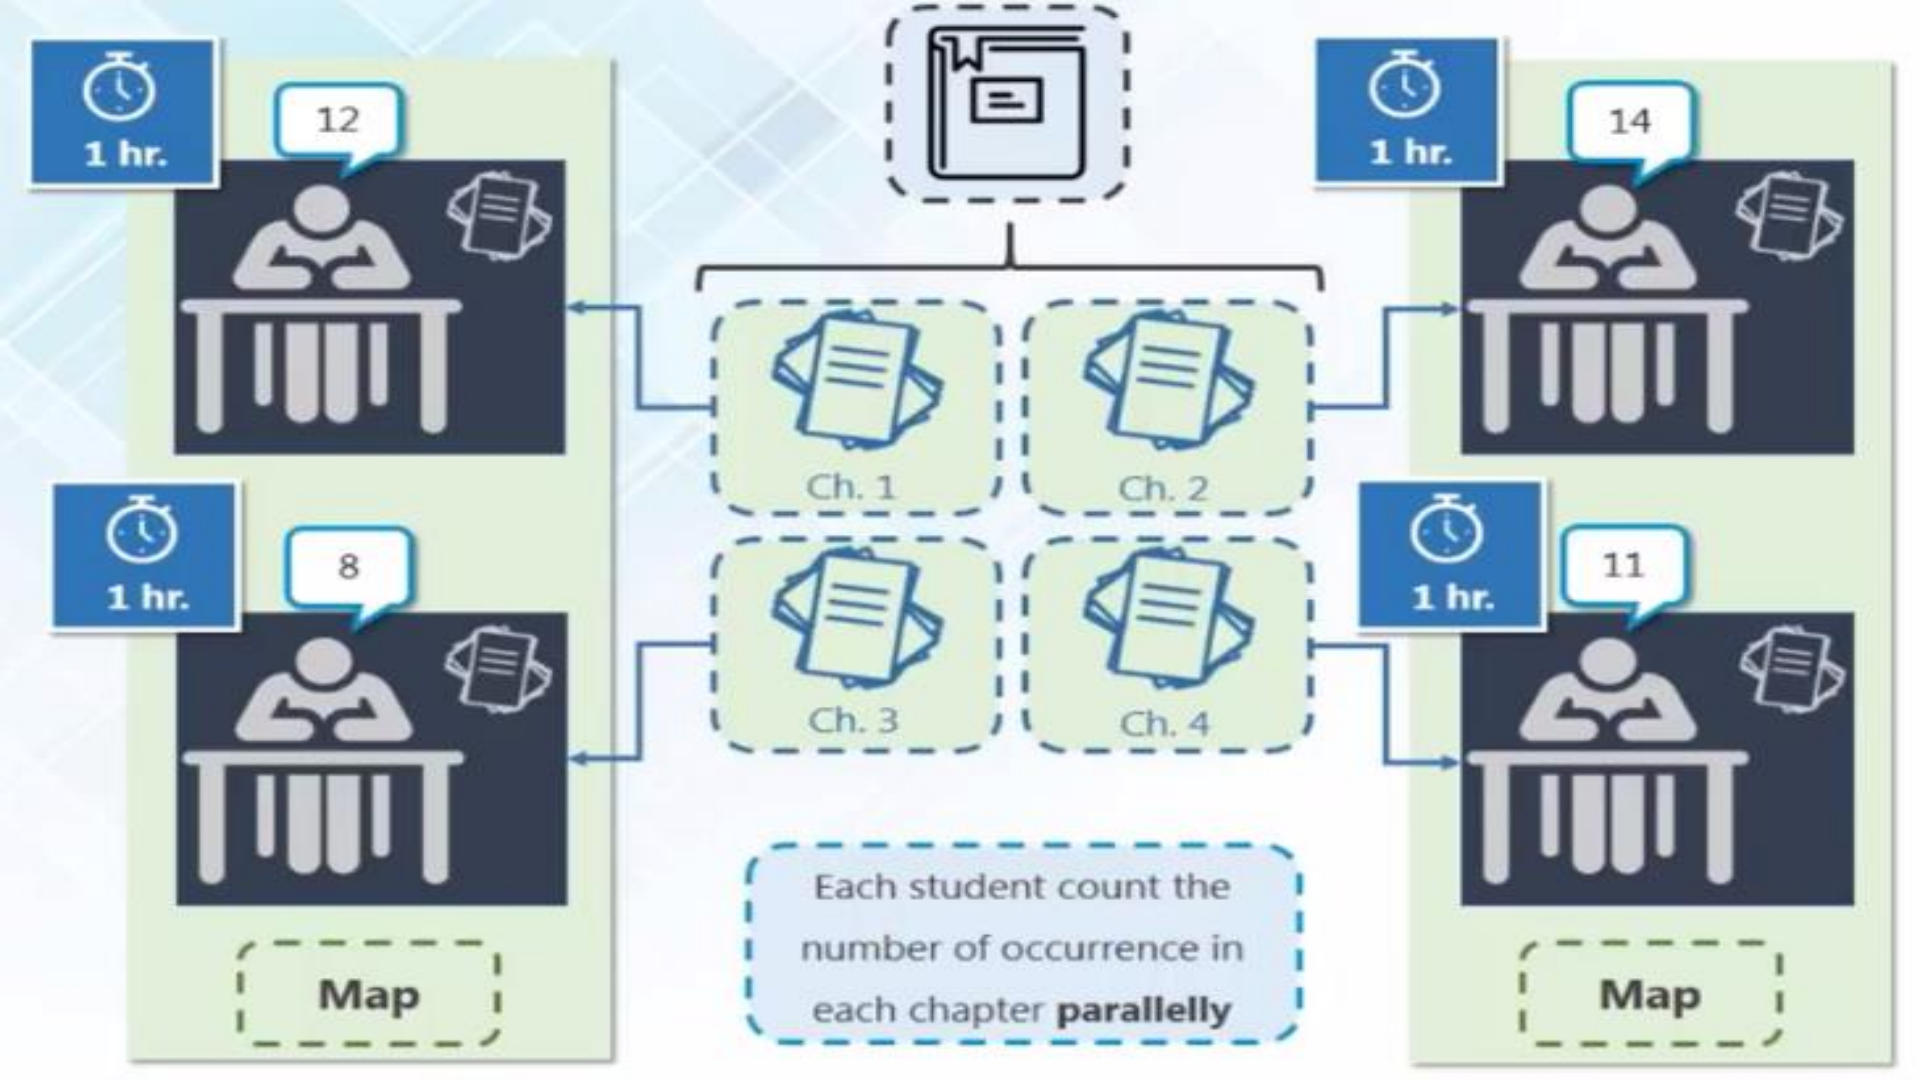
1 hr.
12
1 hr.
14
Ch. 1
Ch. 2
1 hr.
8
Ch. 3
Ch. 4
1 hr.
11
Each student count the number of occurrence in each chapter **parallelly**
Map
Map

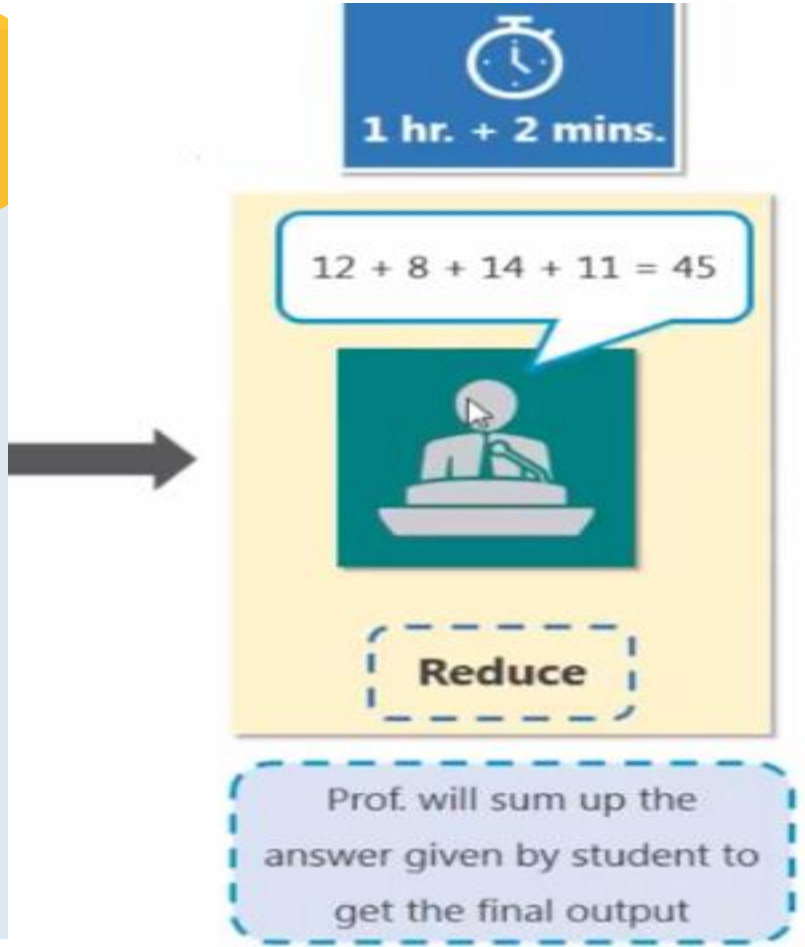
1 hr. + 2 mins.
12 + 8 + 14 + 11 = 45
Reduce
Prof. will sum up the answer given by student to get the final output

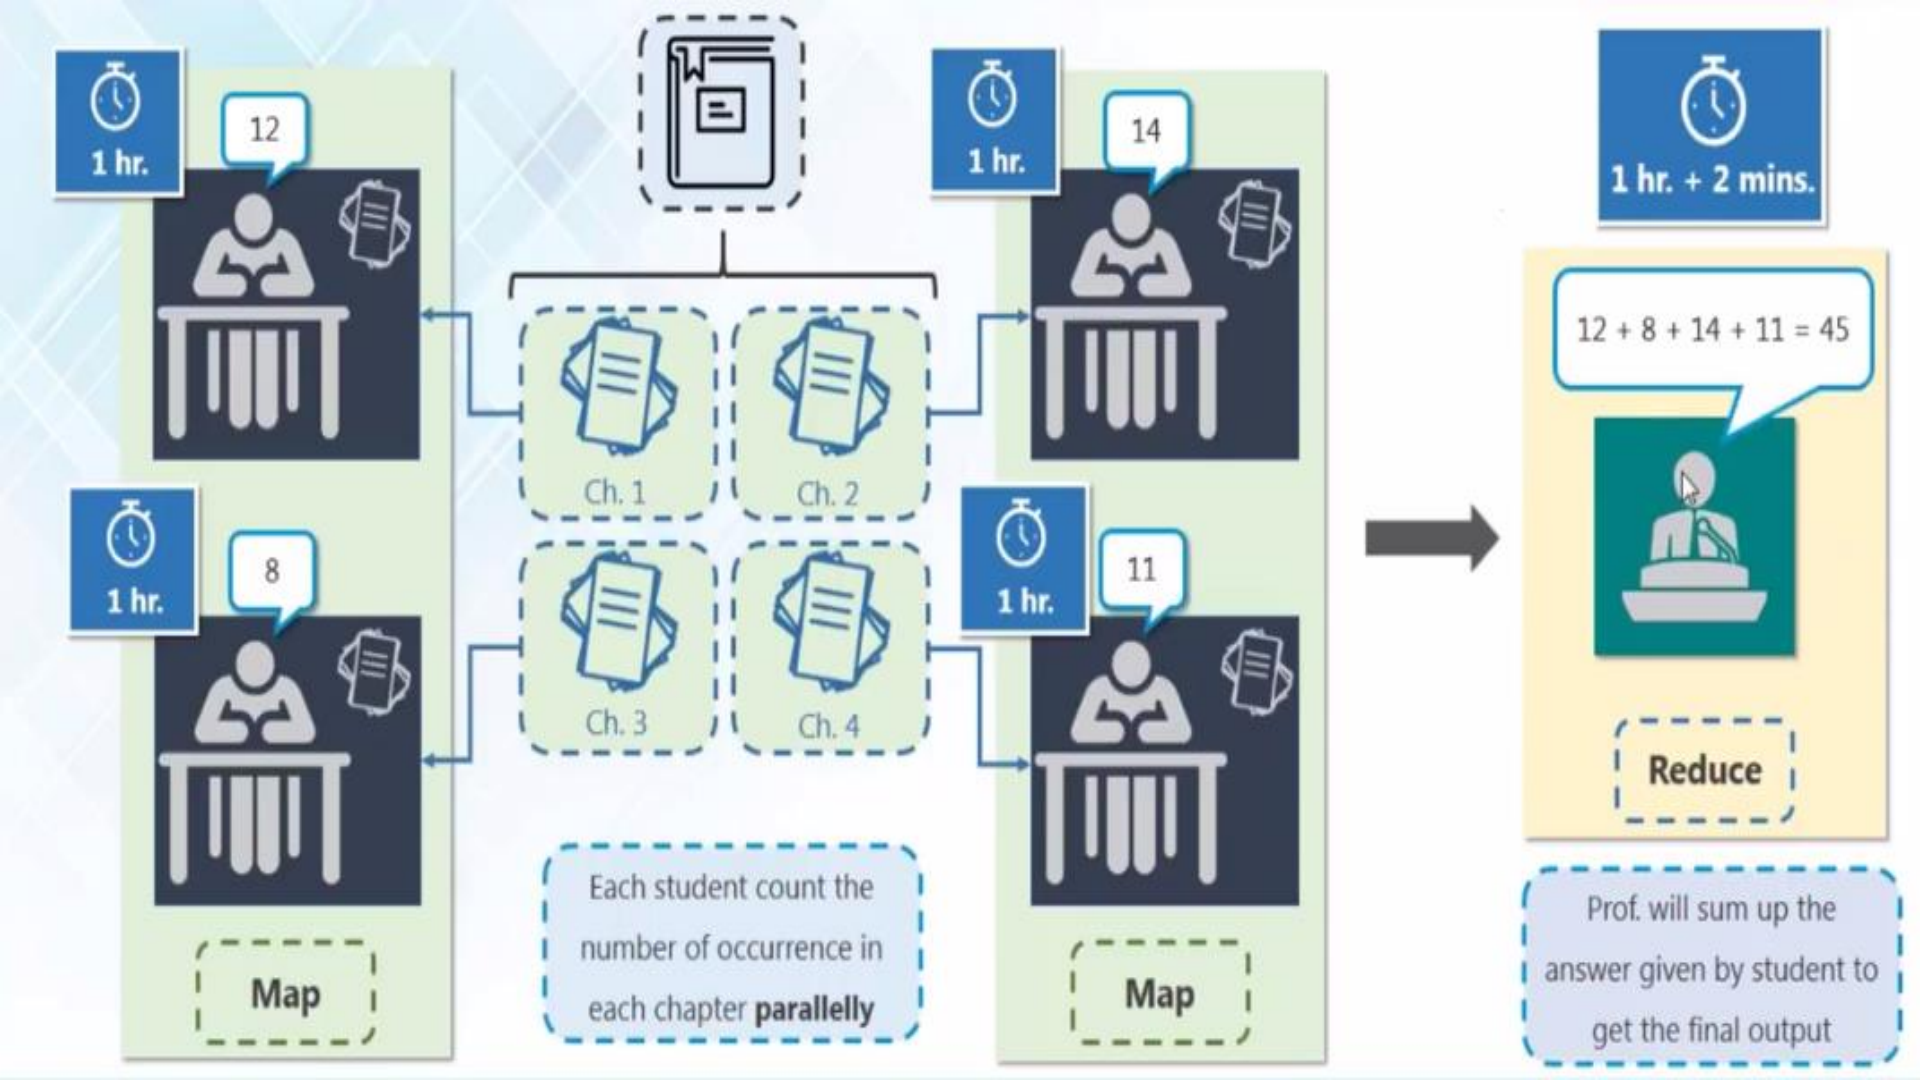
1 hr.
12
1 hr.
14
1 hr. + 2 mins.
Ch. 1
Ch. 2
12 + 8 + 14 + 11 = 45
1 hr.
8
Ch. 3
Ch. 4
1 hr.
11
Reduce
Each student count the number of occurrence in each chapter **parallelly**
Map
Map
Prof. will sum up the answer given by student to get the final output

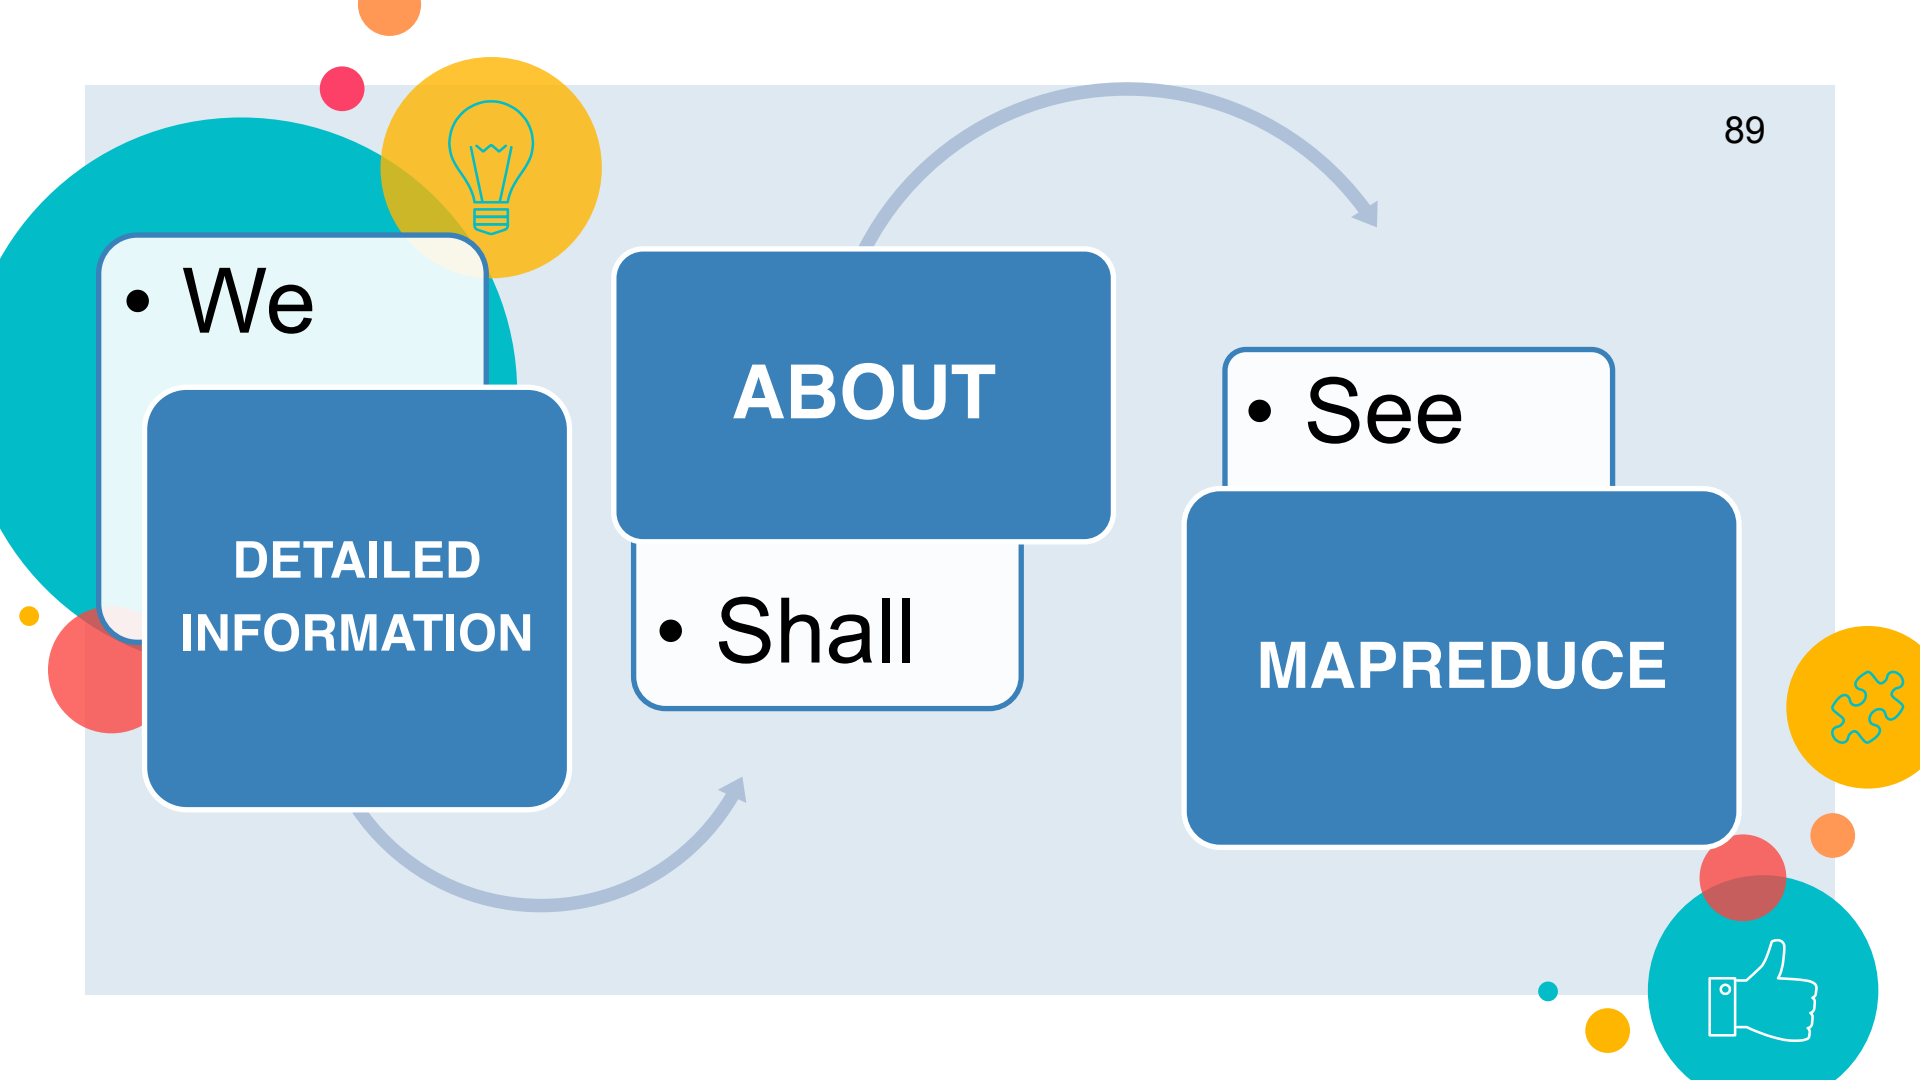
• We
DETAILED INFORMATION
ABOUT
• Shall
• See
MAPREDUCE

MapReduce is a programming framework that allows us to perform distributed and parallel processing on large data sets in a distributed environment

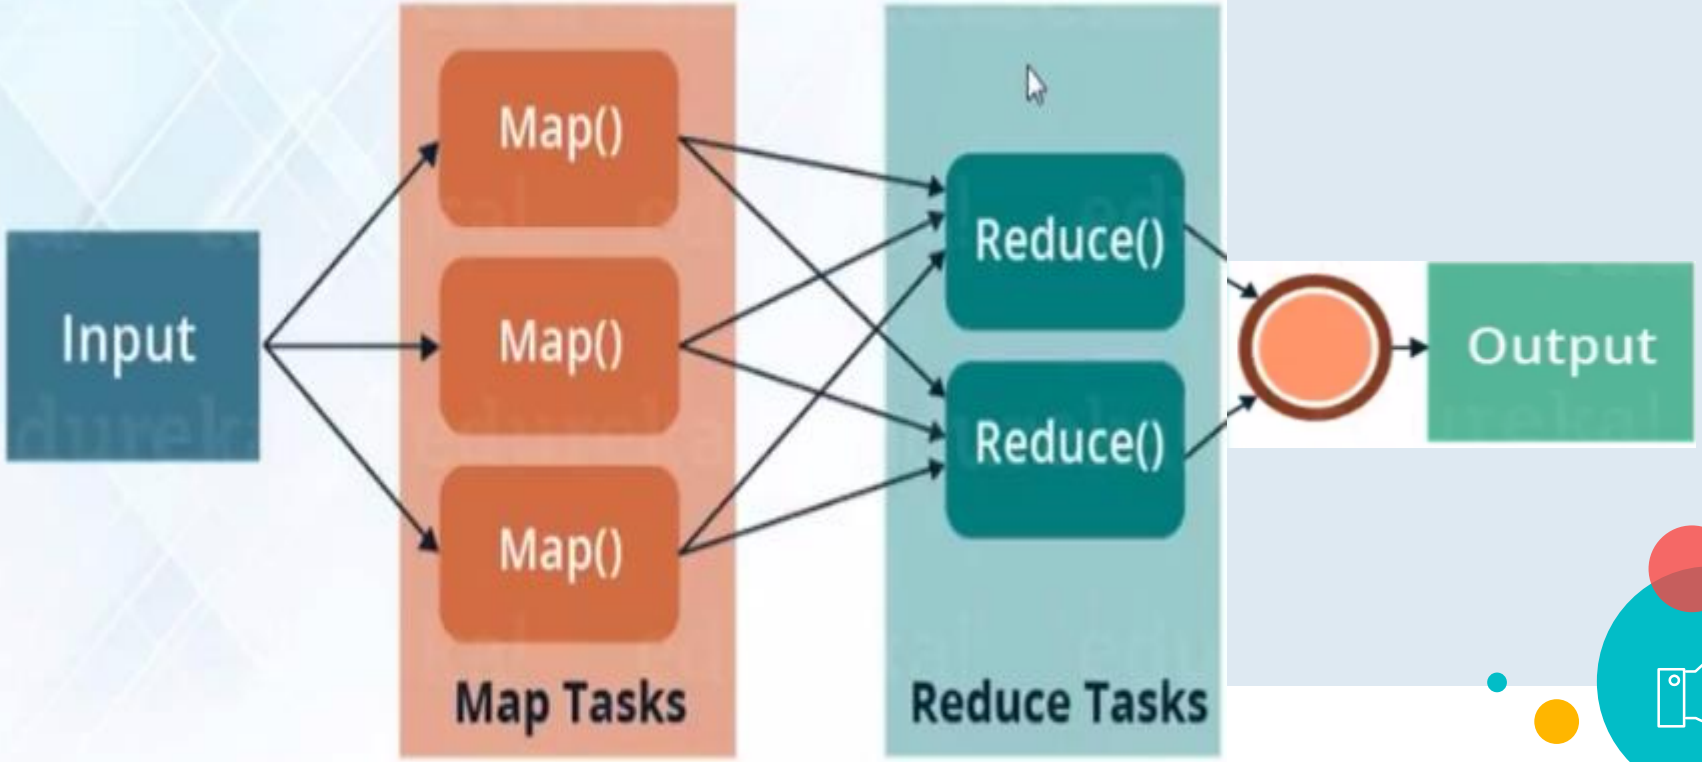

MAP

Reduce

WORD Count Program

The Overall MapReduce Word Count Process

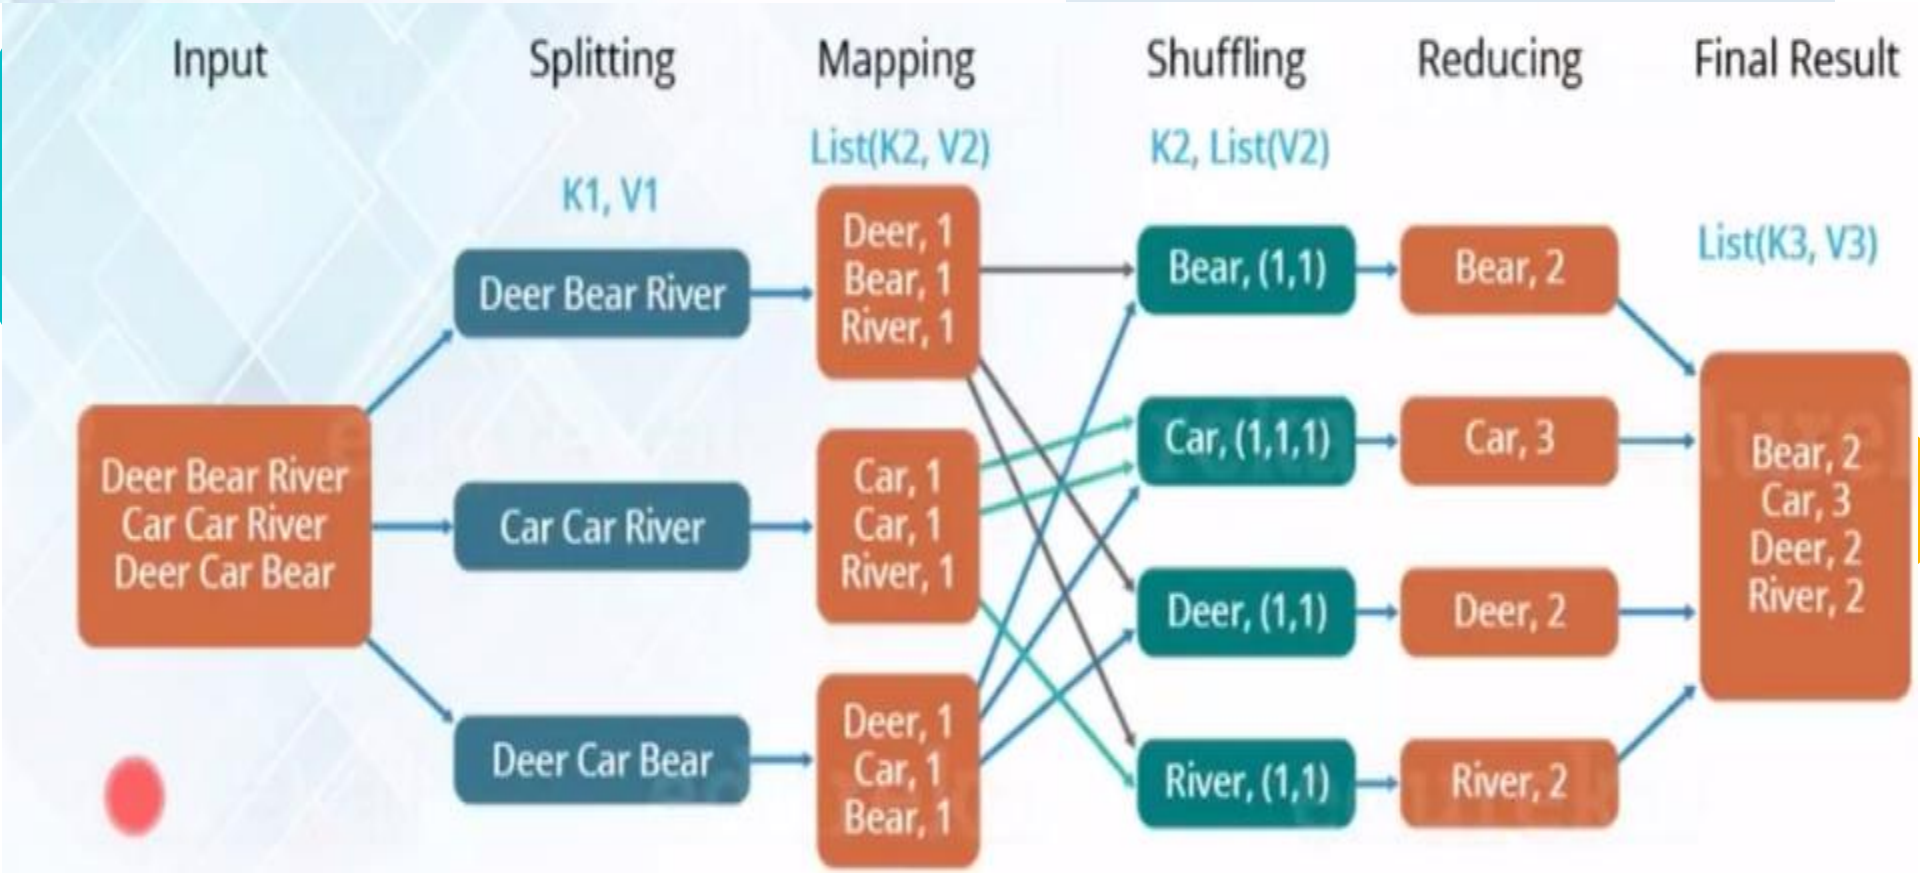

Three Major Parts of MapReduce Program:

1

Mapper Code:

You write the mapper logic over here i.e. how map task will process the data to produce the key-value pair to be aggregated

2

Reducer Code:

You write reducer logic here which combines the intermediate key-value pair generated by Mapper to give the final aggregated output

3

Driver Code

You specify all the job configurations over here like job name, Input path, output path, etc.

YARN Components

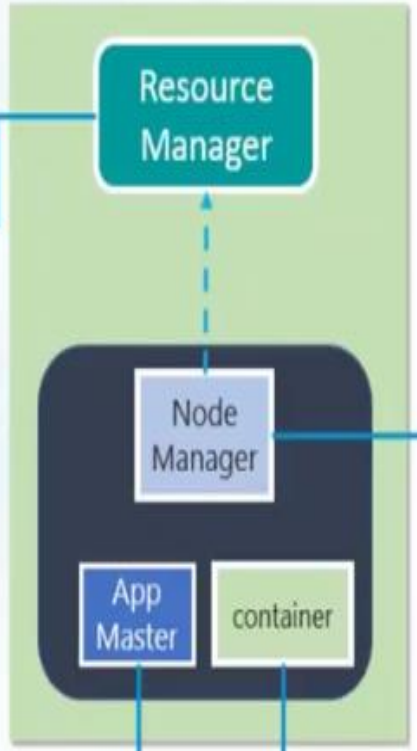

ResourceManager:

- Master daemon that manages all other daemons & accepts job submission
- Allocates first container for the AppMaster

NodeManager:

- Responsible for containers, monitoring their resource usage i.e. (cpu, memory, disk, network) & reports the same to RM

AppMaster:

- One per application
- Coordinates and manages MR Jobs
- Negotiates resources from RM

Container:

- Allocates certain amount of resources (memory, CPU etc.) on a slave node (NM)

MAPREDUCE JOB WORKFLOW

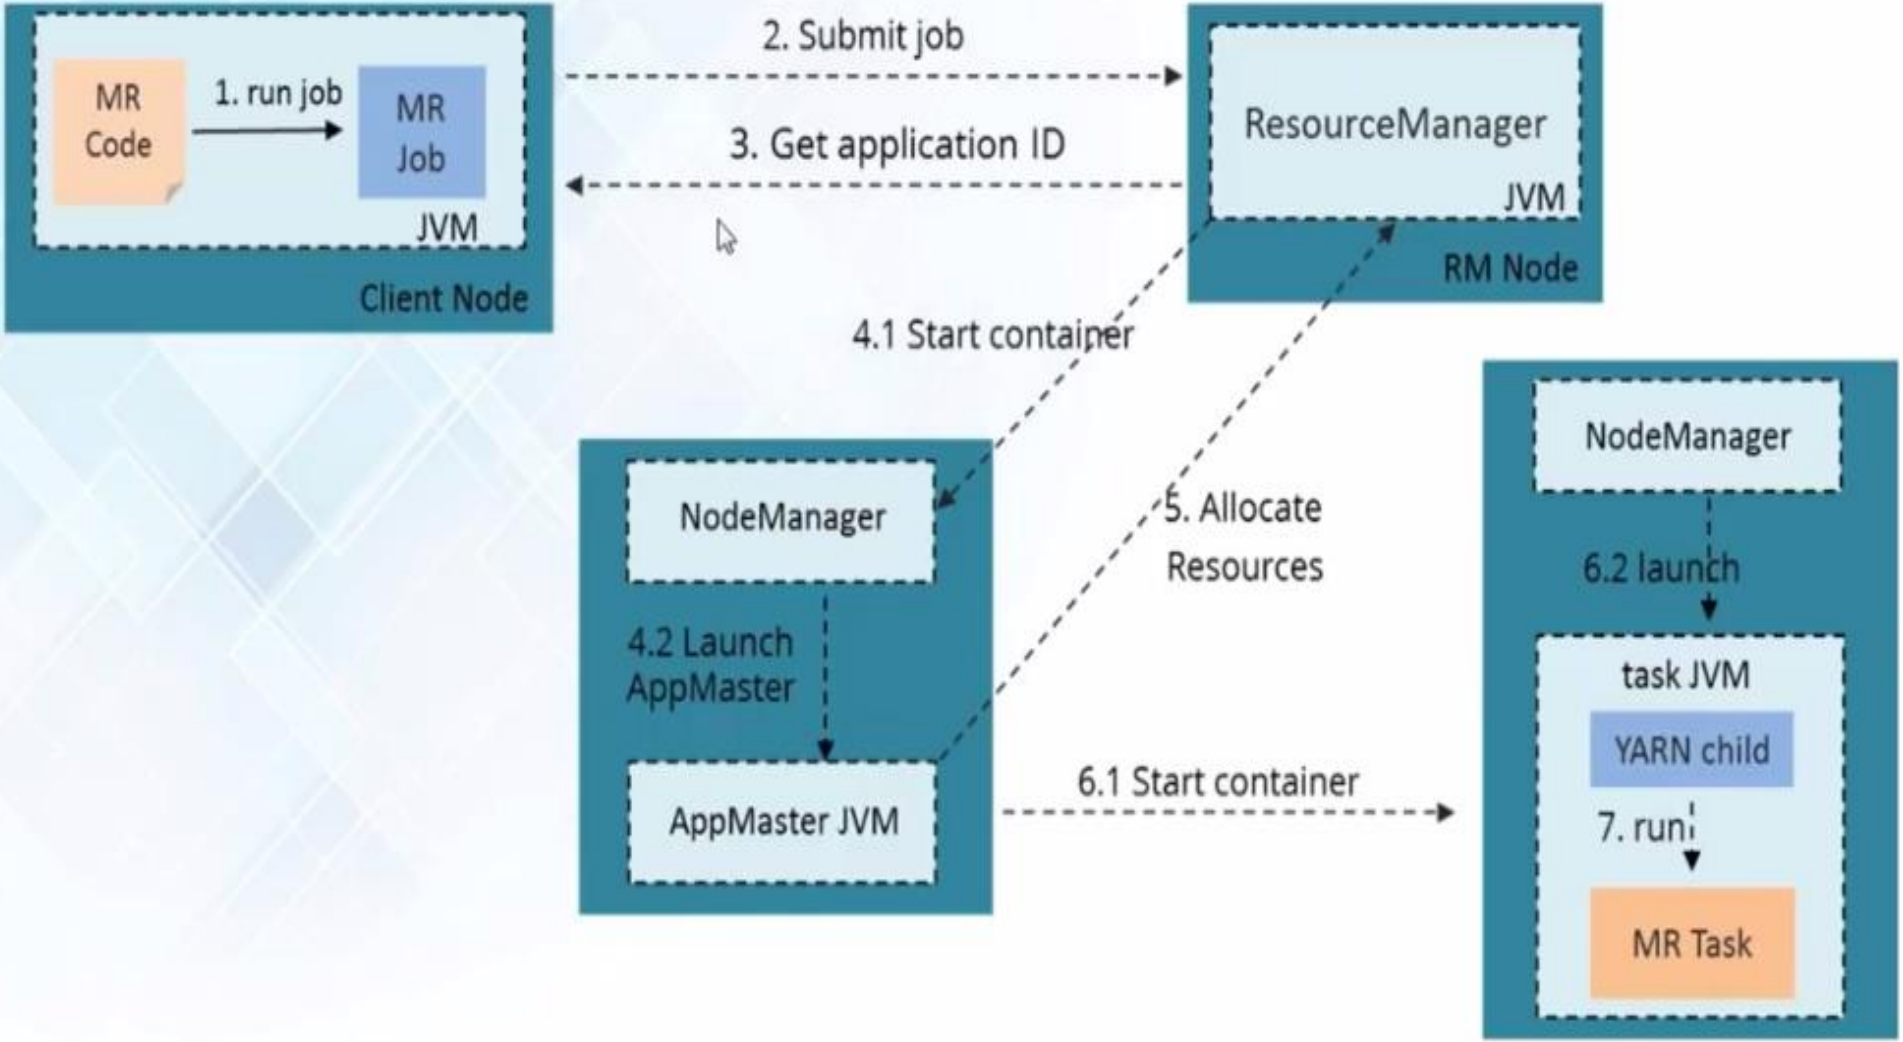
MR Code
1. run job
MR Job
JVM
Client Node
2. Submit job
3. Get application ID
ResourceManager
JVM
RM Node
4.1 Start container
NodeManager
4.2 Launch AppMaster
AppMaster JVM
5. Allocate Resources
6.1 Start container
NodeManager
6.2 launch
task JVM
YARN child
7. run
MR Task