

```

Applications Places System
cloudera@cloudera-vm: ~
File Edit View Search Terminal Help
cloudera@cloudera-vm:~$ sudo jps
[sudo] password for cloudera:
1457 NameNode
1374 JobTracker
1626 TaskTracker
2592 Jps
1801 RunJar
1677 HMaster
1842 DataNode
1538 SecondaryNameNode
cloudera@cloudera-vm:~$

```

1

```

cloudera@cloudera-vm: ~
File Edit View Search Terminal Help
hbase(main):003:0> list
TABLE
t1
1 row(s) in 0.0120 seconds
hbase(main):004:0>

```

2

```

File Edit Format View Help

Create Table

    create 'table_name', 'column_family'

Store Data

    put 'table_name', 'ROW_KEY', 'column_family:column_name', 'value'
    put 'table_name', 'ROW_KEY', 'column_family:column_name_2', 'value'

Get the value from hbase

    (select *)    scan 'table_name'
                  get  'table_name','ROW_KEY'

Update/Modify

    put 'table_name', 'ROW_KEY', 'column_family:column_name', 'value_modified'

Delete data

    delete 'table_name', 'ROW_KEY', 'column_family:column_name'

Drop/alter table

```

3

```

hbase(main):004:0> create 'htest', 'cf'
0 row(s) in 1.1690 seconds

hbase(main):005:0> list
TABLE
htest
t1
2 row(s) in 0.0090 seconds

hbase(main):006:0>

```

4

```

File Edit View Search Terminal Help
hbase(main):004:0> create 'htest' , 'cf'
0 row(s) in 1.1650 seconds

hbase(main):005:0> list
TABLE
htest
t1
2 row(s) in 0.0090 seconds

hbase(main):006:0> put 'htest' , 'r1' , 'cf:c1' , 'v1'
0 row(s) in 0.1020 seconds

hbase(main):007:0> put 'htest' , 'r1' , 'cf:c2' , 'v2'
0 row(s) in 0.0510 seconds

hbase(main):008:0> put 'htest' , 'r1' , 'cf:c3' , 'v3'
0 row(s) in 0.0080 seconds

hbase(main):009:0> scan 'htest'
ROW COLUMN+CELL
r1 column=cf:c1, timestamp=1414254248004, value=v1
r1 column=cf:c2, timestamp=1414254280456, value=v2
r1 column=cf:c3, timestamp=1414254290246, value=v3
1 row(s) in 0.0590 seconds

hbase(main):010:0>

```

5

```

hbase(main):009:0> scan 'htest'
ROW COLUMN+CELL
r1 column=cf:c1, timestamp=1414254248004, value=v1
r1 column=cf:c2, timestamp=1414254280456, value=v2
r1 column=cf:c3, timestamp=1414254290246, value=v3
1 row(s) in 0.0590 seconds

hbase(main):010:0> get 'htest' , 'r1'
COLUMN CELL
cf:c1 timestamp=1414254248004, value=v1
cf:c2 timestamp=1414254280456, value=v2
cf:c3 timestamp=1414254290246, value=v3
3 row(s) in 0.0360 seconds

hbase(main):011:0> put 'htest' , 'r1' , 'cf:c3' , 'v3_updated'
0 row(s) in 0.0220 seconds

hbase(main):012:0> put 'htest' , 'r1' , 'cf:c3' , 'v3_updated_3'
0 row(s) in 0.0110 seconds

```

6

```

hbase(main):009:0> scan 'htest'
ROW COLUMN+CELL
r1 column=cf:c1, timestamp=1414254248004, value=v1
r1 column=cf:c2, timestamp=1414254280456, value=v2
r1 column=cf:c3, timestamp=1414254290246, value=v3
1 row(s) in 0.0590 seconds

hbase(main):010:0> get 'htest' , 'r1'
COLUMN CELL
cf:c1 timestamp=1414254248004, value=v1
cf:c2 timestamp=1414254280456, value=v2
cf:c3 timestamp=1414254290246, value=v3
3 row(s) in 0.0360 seconds

hbase(main):011:0> put 'htest' , 'r1' , 'cf:c3' , 'v3_updated'
0 row(s) in 0.0220 seconds

hbase(main):012:0> put 'htest' , 'r1' , 'cf:c3' , 'v3_updated_3'
0 row(s) in 0.0110 seconds

```

7

```

hbase(main):011:0> put 'htest' , 'r1' , 'cf:c3' , 'v3_updated'
0 row(s) in 0.0220 seconds

hbase(main):012:0> put 'htest' , 'r1' , 'cf:c3' , 'v3_updated_3'
0 row(s) in 0.0110 seconds

hbase(main):013:0> scan 'htest'
ROW COLUMN+CELL
r1 column=cf:c1, timestamp=1414254248004, value=v1
r1 column=cf:c2, timestamp=1414254280456, value=v2
r1 column=cf:c3, timestamp=1414254592531, value=v3_updated_3
1 row(s) in 0.0760 seconds

hbase(main):014:0> get 'htest' , 'r1' , {COLUMN=>'cf:c3',VERSIONS=>3}
COLUMN CELL
cf:c3 timestamp=1414254592531, value=v3_updated_3
cf:c3 timestamp=1414254579862, value=v3_updated
cf:c3 timestamp=1414254290246, value=v3
3 row(s) in 0.0460 seconds

hbase(main):015:0> get 'htest' , 'r1' , {COLUMN=>'cf:c3',VERSIONS=>2}
COLUMN CELL
cf:c3 timestamp=1414254592531, value=v3_updated_3
cf:c3 timestamp=1414254579862, value=v3_updated
2 row(s) in 0.0490 seconds

```

8

```
hbase(main):026:0> put 'htest', 'r1', 'cf:cf:altered:c3', 'v3_updated_6'
0 row(s) in 0.0180 seconds
```

```
hbase(main):027:0> scan 'htest'
```

```
ROW COLUMN+CELL
r1 column=cf:cf:altered:c3, timestamp=1414255410250, value=v3_updated_6
r1 column=cf:cf:c1, timestamp=1414254248004, value=v1
r1 column=cf:cf:c2, timestamp=1414254280456, value=v2
r1 column=cf:cf:c3, timestamp=1414255088377, value=v3_updated_6
1 row(s) in 0.0200 seconds
```

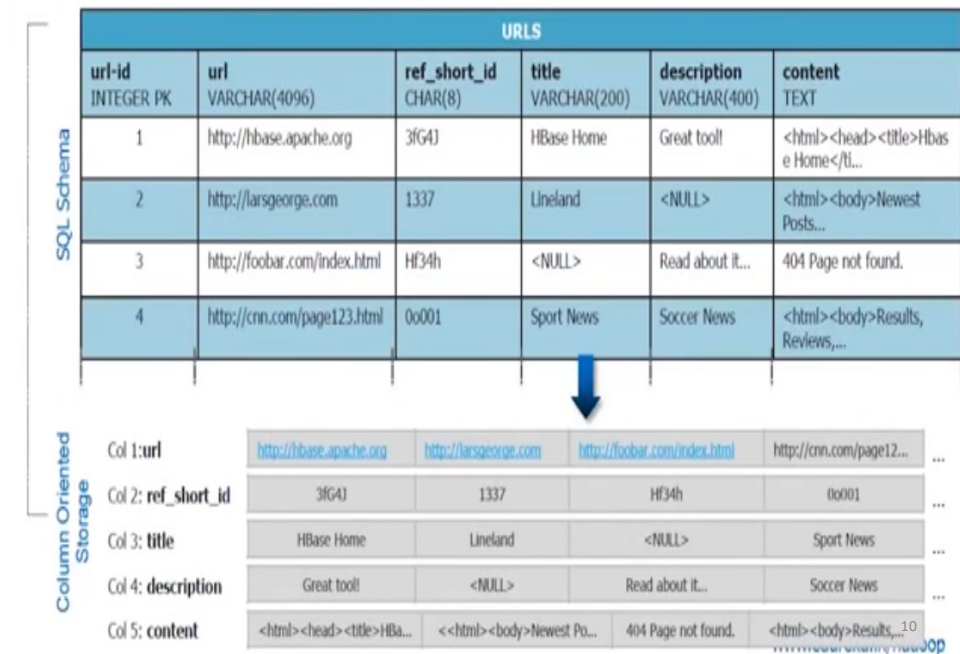
```
hbase(main):028:0> delete 'htest', 'r1', 'cf:cf:c3'
0 row(s) in 0.0140 seconds
```

```
hbase(main):029:0> scan 'htest'
```

```
ROW COLUMN+CELL
r1 column=cf:cf:altered:c3, timestamp=1414255410250, value=v3_updated_6
r1 column=cf:cf:c1, timestamp=1414254248004, value=v1
r1 column=cf:cf:c2, timestamp=1414254280456, value=v2
1 row(s) in 0.0330 seconds
```

9

Row Vs. Column Oriented DBS

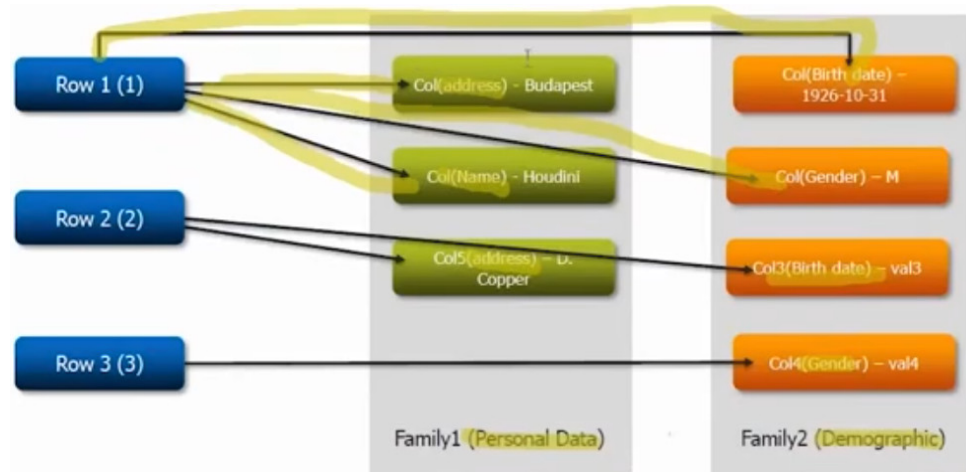


Row Vs. Column Oriented DBS

Row key	Personal_data		demographic	
Persons ID	Name	Address	Birth Date	Gender
1	H. Houdini	Budapest	1926-10-31	M
2	D. Copper		1956-09-16	M
3	Merlin		1136-12-03	F
4	M
500,000,000	F. Cadillac	Nevada	1964-01-07	M

11

Row Vs. Column Oriented DBS



12

How Does it Look Like ?

```

ROW                                COLUMN+CELL
row1                               column=data:1, timestamp=1262873693421, value=value
1
row2                               column=data:2, timestamp=1262873716906, value=value
2
row3                               column=data:3, timestamp=1262873738843, value=value
3
row(s) in 0.0150 seconds
  
```

What it means?

Row Key	Column Family: Column Qualifier	Values
<ul style="list-style-type: none"> Unique for each row Identifies each row 	<ul style="list-style-type: none"> Less number of families gives faster access Families are fixed column qualifiers are not 	<ul style="list-style-type: none"> Various versions of values are maintained Scan shows only recent version

13

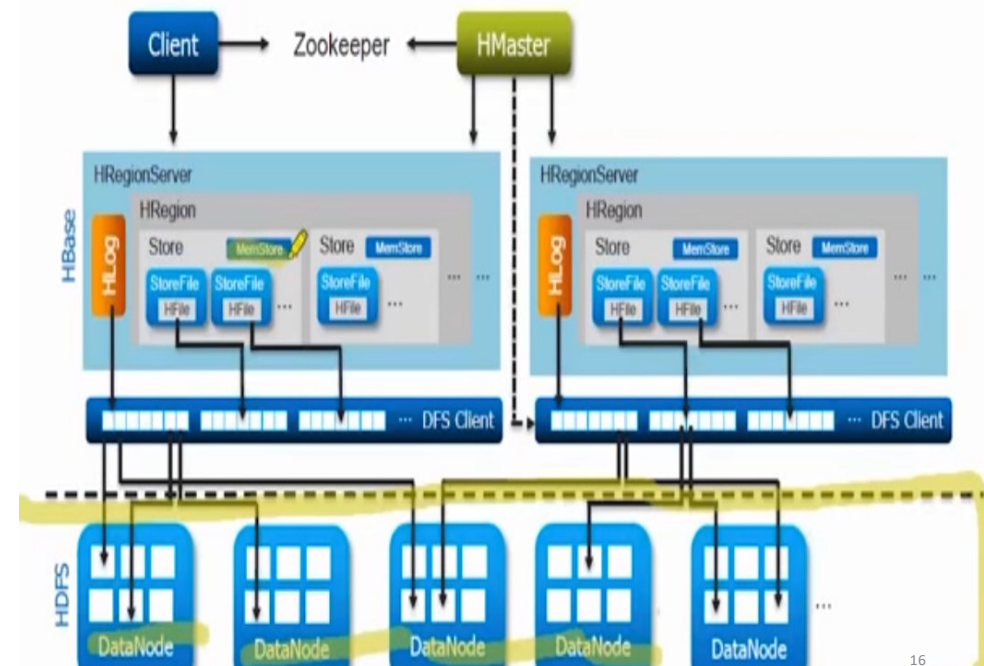


14

HBase Components

- ✓ Table made of regions
- ✓ Region – a range of rows stored together
- ✓ Region servers- serves one or more regions
 - ✓ A region is served by only one region server
- ✓ Master server – daemon responsible for managing HBase cluster
- ✓ HBase stores its data into HDFS
 - ✓ Relies on HDFS's High Availability and fault tolerance

15



16

Cassandra

Apache's Cassandra was chosen by Instagram as the perfect solution to this problem, because of the following features it provides:

- Fully distributed with no single point of failure
- Free and open source with deep developer support
- Linearly Scalable
- Larger-than-memory Datasets
- Built-in-class performance
- Fully Durable
- Integrated Caching and Tuneable Consistency

17

Cassandra

Apache Cassandra is a **free and open-source distributed NoSQL database management system** designed to **handle large amounts of data** across many commodity servers, providing **high availability** with no single point of failure.



18

Cassandra




19

Cassandra

EASY DATA DISTRIBUTION

- Data distribution is very easy
- Provides the *flexibility to distribute data*
- *No master-slave issues* due to peer-to-peer architecture.



ELASTIC SCALABILITY

- Cassandra *scales horizontally*
- During scaling, *Read and Write throughput increase simultaneously*
- There is a *Zero downtime*.

20

Cassandra

HIGH AVAILABILITY & FAULT TOLERANT

- Data is *automatically replicated* to multiple nodes *for fault-tolerance*
- Failed nodes can be replaced with no downtime, due to replication

EFFICIENT WRITES

- Was designed to *run on cheap commodity hardware*
- It performs *blazingly fast writes*
- Can store hundreds of terabytes of data, without sacrificing the read efficiency


CQL

CQL

- Cassandra introduced the *Cassandra Query Language(CQL)*
- CQL 3 is the default and the primary interface into the Cassandra DBMS

21

Cassandra Features of Cassandra




TUNABLE CONSISTENCY

- Provides a means for *tuning the level of consistency required*
- Consistency makes sure that the *client is approved as soon as the cluster accepts the write*

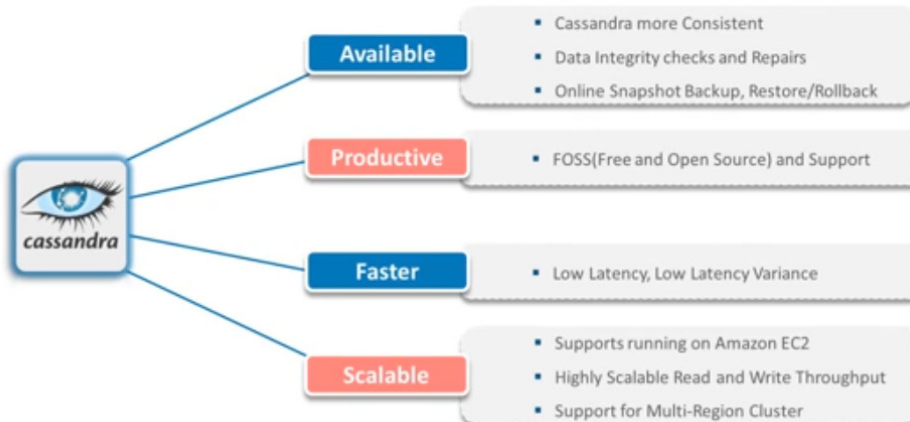
FLEXIBLE DATA STORAGE

- Cassandra accommodates *all possible data formats*
- Structured, Semi-structured, and Unstructured*
- Can *dynamically accommodate changes* to data structures



Cassandra

Reasons to Chose Cassandra?



23

Cassandra

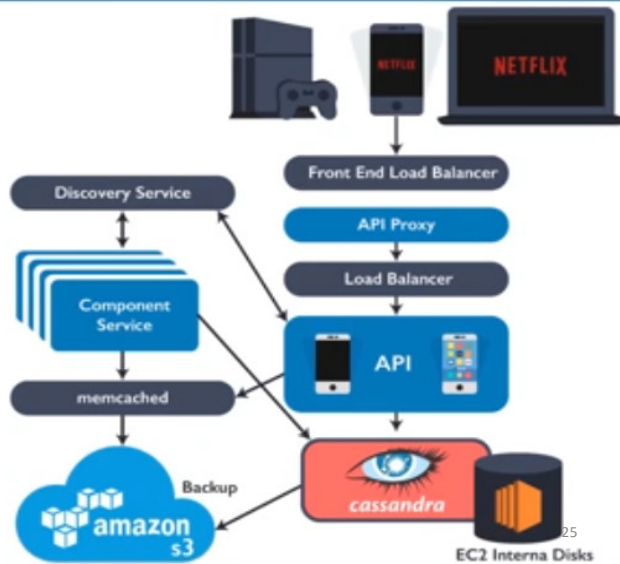
How Extensively Netflix Uses Cassandra



24

Cassandra

Cassandra in Netflix's Architecture




Cassandra

CAP Theorem

- Eric Brewer posited his CAP theorem in 2000, which states within a large-scale distributed data system, you can only have two out of the following three guarantees:

- Consistency
- Availability
- Partition tolerance




26

Cassandra

CONSISTENCY

Read operation will return the value of the most recent write operation, causing all nodes to return the same data



AVAILABILITY

Every request gets a response on Success/Failure
System remains operational 100%. Every client gets a response, regardless of any individual node in the system

PARTITION TOLERANCE

System continues to work regardless of partial failure
Sustain any amount of network failure, which does not result in failure of the entire network



27